

FlowPsi: An ideal gas flow solver – The User Guide.

Edward A. Luke, Xiaoling Tong, and Rex Chamberlain

January 7, 2023

Contents

1	Introduction	3
2	Units	3
3	Problem Input	3
3.1	Problem Description File	3
3.2	Boundary Condition Specifications	12
3.2.1	Notes on Vector Value Inputs	15
3.2.2	Notes on Inflow Velocity Specifications	16
3.2.3	Notes on Setting Viscous Boundary Velocity	16
3.2.4	Integrated Forces and Moments	17
4	Relative Motion Problems using <i>FlowPsi</i>	18
4.1	Specifying motion in a relative frame	19
4.2	Hole Cutting	19
5	<i>FlowPsi</i> Pre- and Post-processing Utilities	19
5.1	<i>FlowPsi</i> Pre-processing Utilities	19
5.1.1	ugrid2vog	19
5.1.2	xdr2vog	20
5.1.3	vogmerge	20
5.1.4	cobalt2vog	20
5.1.5	plot3d2vog	20
5.1.6	fluent2vog	20
5.1.7	ccm2vog	21
5.1.8	vogcheck	21
5.2	<i>FlowPsi</i> Post-processing utilities	21
5.2.1	Probe Output	21
5.2.2	Plot Format Converter	22
5.3	Surface Data Output Options	23
5.4	Mean and Variance Output	24
5.5	The Grid Motion Module	25
5.5.1	Theoretical Background	25
5.5.2	Optimizations	26
5.5.3	Grid Motion User Controls	29
A	Governing Equations	32
A.1	Thermodynamic Model	32

FlowPsi User Guide (Version 1.0)	2
B Formulation for Axisymmetric Flows with Swirl	32
C Turbulence Models	35
C.1 Spalart-Allmaras Model	35
C.2 Compressible Form of Spalart-Allmaras	35
C.3 Baseline Model (BSL)	36
C.4 Shear Stress Transport Model (SST)	38
C.5 The Wilcox (1998) $k - \omega$ Turbulence Model	38
C.6 Hybrid RANS/LES Model (LES)	39
C.7 Turbulence Compressibility Corrections	41
C.8 The Dynamic Hybrid RANS-LES (DHRL) turbulence model	41
C.8.1 Current DHRL Implementation	42
C.8.2 DHRL Inputs	42
C.8.3 Optional inputs for DHRL include:	43
C.8.4 DHRL Outputs	43
C.8.5 Recommended DHRL Start Up Procedure	43
C.8.6 Optimal settings for MILES	44
D Numerical Formulation	44
D.1 Numerical Approximations to Spatial Integrals	44
D.2 Geometric Integrations	46
D.3 Inviscid Flux Schemes	47
D.4 Preconditioned HLLC flux option	47
D.4.1 Background on preconditioning	47
D.4.2 Skew Symmetric low dissipation flux schemes	48
D.4.3 Unstructured Mesh Hybridization	49
D.4.4 Hybrid Flux Scheme	50
D.5 Time Integration	52
D.5.1 Local Time-Stepping Scheme	53
D.6 Scalar Transport Equations	53

1 Introduction

FlowPsi is an open source CFD solver that is developed utilizing the Loci framework[1, 2] which provides for automatic parallelization and logical consistency checks simplifying the design of highly parallel numerical solvers for partial differential equations.

2 Units

All input into the solver is assumed to be in the international MKS system. For example, length is measured in meters, mass in kilograms, time in seconds, pressure in Pascals, temperature in Kelvin, heat in Joules, and so on. Unless otherwise noted this standard is maintained in all *FlowPsi* input and output files. Note that *FlowPsi* provides unit conversions on most inputs such that a user can input the units that are most comfortable. Output files are always in MKS units, however.

3 Problem Input

The input to *FlowPsi* consists of a problem description file containing variable definitions and postfixed with `.vars`. The file describes the boundary conditions, initial conditions, and directives to the numerical solver. The grid file is stored in a Loci specific “volume grid” format that has a `.vog` postfix. A variety of mesh conversion utilities are provided as part of the Loci framework to convert the grid to the format used by *FlowPsi*.

3.1 Problem Description File

The problem description file (with the postfix `.vars`) is used to describe the grid, boundary conditions, initial conditions, and numerical solution parameters. The occurrence of the characters `//` anywhere in the file indicates a comment to the end of that line. The file consists of a set of variable definitions enclosed in braces. Each variable is listed at the beginning of a line followed by a colon and then its definition. An example of an input file named `input.vars` is as follows:

```
{
boundary_conditions: <
    body=viscousWall(adiabatic),
    base=viscousWall(adiabatic),
    BC_7=fixedMass(mdot = 0.0027 kg/s, T0=300 K) //air injection
    BC_5=symmetry,
    BC_6=symmetry,
    BC_10=reflecting,
    outflow=extrapolate,
    inflow=supersonicInflow(p=21571.6 Pa, T=149.6 K, M=2.2)>

// initial conditions
initialConditions: < p=21571.6, T=149.6, M=0.0>
p0: 21571.6 //gage pressure reference

gridCoordinates: axisymmetric // add axisymmetric frame of reference
flowRegime: laminar
timeStepMode: steady
```

```

plot_freq: 100
plot_modulo: 100
restart_freq:100
restart_modulo: 100
stop_iter: 5000

```

```

limiter: venkatakrishnan
K1: 1.0

```

```

urelax: 0.20
dtmax: 1.0e-3
}

```

You can find out about all of the available input variables to *FlowPsi* by using the “-inputs” option to the *FlowPsi* executable. These variables are also defined as follows:

1. boundary_conditions

This variable list is used to associate boundary surface tags in the grid file with a boundary condition specification. For example:

```

boundary_conditions: <
    Body=viscousWall(adiabatic),
    SymmetryPlane=symmetry,
    OutflowPlane=outflow(p=5005.1342psi),
    InflowPlane=supersonicInflow(p=1atm, T=500K, M=1.2) >

```

Note: The last entry of this specification does not have a comma! Also, make sure that in the grid generator program (Gridgen, for example) you use different flags for each of the boundary surfaces that you will want to process separately (*e.g.*, integrate quantities, visualize, etc.). Note, the terms **Body**, **SymmetryPlane**, etc. are symbolic names given to these boundary surfaces by the grid generator. If the symbolic names are not available when the grid file is converted to the **vog** format, then these names will be related to the original integer tag given to the surface. For example, if a surface had an integer tag of 5, then the name **BC.5** would be given to the surface if no symbolic name was provided.

2. initialConditions

These variables define the uniform initial conditions throughout the domain (Note: the initial conditions are automatically read from a file when restarting the solver, and it is also possible to specify an interpolated initial conditions file - see below). The initial conditions are specified by a set of flags and options as described in Table 1. The variable **initialConditions** is used to define the uniform initial values. To describe initial conditions with different values in different regions, see the next item. For specification of the initial conditions, only two thermodynamic inputs are required from the choice of **P**, **T**, or **rho** and that either **u** or **M** is also needed to define velocity. If a turbulence model is loaded, then additional variables may be specified in the initial conditions to override default values (see the turbulence model section for the more initial condition information).

An example for specifying the initial conditions in the **.vars** file is as follows:

```

initialConditions: <p=10 atm, T=550 R, M=0.1>

```

Variable	description	units	notes
<code>rho</code>	density	kg/m^3	specified with either <code>p</code> or <code>T</code>
<code>p</code>	pressure	Pa	specified with either <code>T</code> or <code>rho</code>
<code>T</code>	temperature	K	specified with either <code>p</code> or <code>rho</code>
<code>u</code>	velocity	m/s	exclusive with <code>M</code>
<code>M</code>	Mach number	–	exclusive with <code>u</code>

Table 1: variables describing a fluid state input

3. `initialConditionRegions`

This option is used to specify initial conditions with regions of constant properties. Regions can be defined using geometric primitives such as boxes, spheres, cylinders, and so on. In the `initialConditionRegions` specification, there are two required options: the definition of the default initial condition state, and the definition of a list of regions (defined by geometric primitives). Each geometric region can refer to other defined states. For example, to set up a spherical blast problem, one can define the initial conditions using `initialConditionRegions` as follows:

```
initialConditionRegions: <
default = state(p=1atm,T=298K,u=0),
high_pressure = state(p=1000atm,T=5000K,u=0),
regions = [ inSphere(radius=1m,center=[0,0,0],composition=high_pressure) ]
>
```

The state function contains a specification of the state of the fluid system and can contain any of the initial conditions parameters. The regions variable describes a list of geometric regions. Each geometric region has a corresponding composition which is in turn described by an assignment of a state function. The last specified region has priority over all previous regions. Thus one could make a cone with a rounded outer boundary by combining a conic region and a sphere as follows:

```
initialConditionRegions: < default = state(p=1atm,T=298K,u=0),
high_pressure1 = state(p=1000atm,T=5000K,u=0),
high_pressure2 = state(p=500atm,T=1000K,u=0),
regions = [ inSphere(radius=1m,center=[1,0,0],composition=high_pressure2),
inConic(r1=0,r2=1,p1=[-5,0,0],p2=[1,0,0],composition=high_pressure1) ]
>
```

The geometric primitives currently supported are as follows:

- **`inBox`**
`inBox` takes two arguments, `p1` and `p2`, which are two diagonally opposing corners that define the box. The geometry is defined as any point that lies inside of this rectangular axis aligned region.
- **`inSphere`**
`inSphere` takes two arguments, `radius` and `center`, which define the sphere size and location, respectively.
- **`inCylinder`**
`inCylinder` takes three arguments, `radius`, `p1`, and `p2`. The endpoints of the axis are defined by the points `p1` and `p2`.

- **inCone**
inCone takes four arguments. Like **inCylinder**, it uses two endpoints of the axis, while **r1** and **r2** define the respective radii at each endpoint.
- **leftPlane**
leftPlane takes two arguments, **point** and **normal**. **point** defines a point on the plane, while **normal** gives a vector normal to the plane. All points on the opposing side of the plane (away from the normal) are included in this geometry.

4. **inviscidFlux**

Select flux formulation for the inviscid flux in the solver. The selection can be “hllc”, “kec”, or “ssf”. The “hllc” is the default upwind scheme. The “kec” is a second order low dissipation flux, and “ssf” is based on a fourth order skew symmetric scheme. Select the “ssf” for LES style simulations.

5. **flowRegime**

This sets the type of flow model we will be using. This will be set to either **inviscid**, **laminar**, or **turbulent**. If the setting is **turbulent** then a turbulence model will need to be loaded using a **loadModule** directive at the top of the vars file.

6. **timeStepMode**

This variable selects the time integration mode. It is set to either **steady** or **unsteady**. In **steady** mode, **flowPsi** will utilize a backward Euler time integration method with local timestepping enabled. For time-accurate simulations select **unsteady** to disable local timestepping and switch to a second order temporal integration.

7. **timePressureBiasControl**

This variable changes how pressure is interpolated in the unsteady mode. First order biasing is required to improve stability in when using low dissipation fluxes such as ‘kec’ or ‘ssf’ on low speed problems. Setting this to 1.0 will change the temporal biasing to first order. The default value of is 0.0 which will only turn on pressure biasing for cells that exceed a CFL of 1.0.

8. **Minf**

This variable is used to activate preconditioning for low speed flows. Set this variable to the free stream Mach number for the simulation. This variable is used to set bound the minimum preconditioning parameter in stagnation flows to avoid instability. By default this variable is set to 1 which disables low speed preconditioning.

9. **print_freq**

This variable describes the frequency at which CFL numbers and integrated boundary condition data will be displayed on standard output. For example, if this variable is set to 100, then output will be generated for every time step evenly divisible by 100.

10. **restart_freq**

This variable describes the frequency at which restart files are output. Restart files are output in the **restart/#** directory, where **#** is the current iteration number (modulo **restart_modulo**, see next item). This directory contains many files that define the entire restart state. If you want to prevent a restart state from being overwritten in subsequent runs, copy this entire directory to a safe location.

11. `restart_modulo`

This variable controls the naming of the output file. The iteration number modulo this value will be used when outputting restart files (overwriting files if necessary). This can save on disk space by limiting the maximum number of files that may exist at any given time. It is recommended to set it to at least twice the restart frequency to guarantee that the restart file won't become corrupted if the program terminates while writing a restart file.

12. `plot_freq`

Outputs values that are used for generating visualizations and plots. The `extract` utility can be used to convert these files into formats compatible with other post-processing visualization programs such as *FieldView*, *EnSight*, or *Tecplot*. (See the Plot Format Converter in Section 5.2.2).

13. `plot_modulo`

Similar to `restart modulo`, except applies to plot files.

14. `boundary_plot_freq`

The frequency that boundary information will be plotted. Note, boundary plot files will be written at both `plot_freq` iterations and `boundary_plot_freq` iterations.

15. `plot_output`

Provides a comma delimited list of additional variables to output along with the default variables.

16. `plot_output_exclusive`

Provides a comma delimited list of additional variables to output. If this variable is present, then the default output variables will not be written unless they are included in this list. A list of additional variables is given in table 3. ??.

17. `stop_iter`

This variable defines the iteration at which the simulation is terminated.

18. `maximumRunTime`

This variable defines the maximum running time for the simulation. When this time is exceeded the code will automatically save restart files and terminate. The default units is seconds, but the user can specify any time units. For example, if one wanted to make sure the job stopped before 24 hours expired, you would put this in the `.vars` file:

```
maximumRunTime: 23.5 hour
```

19. `newtonMaxIter`

This variable defines the number of Newton iterations performed for each time-step. Typically 1 for steady-state calculations, and ≥ 3 for unsteady calculations.

20. `fluidLinearSolver`

This variable selects the fluid linear system solver to use to solve the fluid and turbulence models. The default setting is `sgs` which uses the built in symmetric Gauss-Seidel solver. The setting `fsgs` uses a fast, cache optimized, symmetric Gauss-Seidel solver. It is faster than `sgs` but uses more memory. A robust and fast solver that is recommended is the

line symmetric Gauss-Seidel solver which is selected by setting this variable to `lsgs`. This solver divides the matrix into lines and solves each line with a tri-diagonal solver. The solver sweeps through lines similar to the symmetric Gauss-Seidel. For the line solver, some under-relaxation is employed to stabilize the solver. In general, this solver is very robust and useful for high Reynolds number or low speed problems. The most robust solver is the one provided by PetSC (if it is installed). This solver is selected by setting this variable to `petsc`.

21. `gauss_seidel_iter`

This variable defines the number of Gauss-Seidel iterations that will be performed when solving the linear system using `sgs`. When this variable is set to zero, the linear solver reduces to a block Jacobi iteration. Typical values are between 5-20. Usually a value of 5-8 is useful for parallel simulations.

22. `LDS_compFactor`

When using the low dissipation fluxes 'kec' or 'ssf' it is necessary to switch to the upwind scheme around compressible features such as shocks. This input sets the sensitivity of the factor to switch to upwind when compressible flow conditions are detected. The default value for this is 5.0. This value is typically set somewhere between 1-100. Set it to zero if these compressible features should not be encountered.

23. `LDS_MILES_filter`

Select how much third order upwinding to use as a MILES filter for the fourth order skew symmetric flux. The default setting for this is 10%.

24. `LDS_useUpwind`

This input will set the lower limit for the upwind blending. Setting this to a small positive value such as 0.2 can be used for a MILES type of LES subgrid model. The default value is 0.0.

25. `LDS_nonSymmetricCoeff`

For unstructured meshes some upwinding is needed to stabilize the scheme. Heuristic methods are used to increase the upwinding blending parameter when mesh cells depart from symmetric stencils. The larger this factor the more robust the scheme will be to poor mesh quality at the expense of increased dissipation. The default value is 2.

26. `LDS_nonSymmetricFactor`

For unstructured meshes this controls the rate at the scheme adds upwinding at small angles. A larger number will increase robustness at the expense of added dissipation. the default value is 1.

27. `LSGSMaxIter`

The maximum number of line Gauss-Seidel steps allowed per solve. The default is 10.

28. `LSGSRelTol`

The residual relative error tolerance for the linear system solver to conclude iteration. The default value is 0.1. Note, that in the *FlowPsi* solver the linear system solved is only an approximation. Usually after some amount of residual convergence it is more productive to proceed to the next Newton iteration in order to achieve overall convergence of the PDE. Over-converging the linear system solver can be non-productive as the linear system itself contains approximation errors. Usually a factor of 10 reduction in the linear system residual

is all that is needed for the Newton iterations to proceed productively, but some problems will benefit from a more robust convergence of the approximate linearization. For difficult problems there may be some productivity to adjusting this parameter to a smaller value (while also increasing the `LSGSMaxIter` variable).

29. `LSGSAbsTol`

The absolute error tolerance for the residual. The default is 1e-10.

30. `LSGSRelaxation`

The under-relaxation parameter for the LSGS solver. Usually the line Symmetric Gauss Seidel is unstable when solving for large timesteps. An under-relaxation parameter of 0.3 (the default) provides stability and rapid convergence for most problems. For problems where the LSGS solver diverges, a smaller setting to this value may allow convergence. However, setting this value too small may prevent convergence in any reasonable number of iterations.

31. `limiter`

This variable describes the limiter that will be used to limit the MUSCL extrapolation scheme. The current choices are `venkatakrishnan` (default), `barth`, `none`, and `zero`. The Venkatakrishnan limiter also has a thresholding parameter that assists in improving accuracy and convergence. The `none` option disables the limiter, and this can only be used on sub-critical flows that do not contain discontinuities. The `zero` option disables higher order extrapolation (first order solution). Generally, the Barth limiter provides robustness while the Venkatakrishnan limiter provides improved accuracy and convergence characteristics.

32. `K1`

This is the thresholding parameter for the Venkatakrishnan limiter. It is a problem dependent parameter that can range from 0 to 100. When `K1` is zero, the limiter is no longer thresholded. If `K1` is set to a large value, limiting is effectively turned off everywhere. At an ideal setting, this parameter allows the limiter to be disabled in smooth regions permitting higher accuracy while still being applied in regions of discontinuities. By default, `K1` is set to 1. The default setting should be sufficient for most uses.

33. `dtmax`

This variable is the global time step parameter.

34. `cflmax`

When this variable is set to a non-zero value, it enables local time-stepping. Once local time-stepping is enabled, each cell's time-step is determined by the smallest of the following candidate time-steps: 1) `dtmax`, 2) the time-step determined from `cflmax`, and 3) a time-step determined from the `urelax` (see item 35) parameter.

35. `urelax`

This parameter limits local time-steps (when `cflmax > 0`) to a time-step that would produce no more than a `urelax` factor change in density, pressure, and temperature. When set to a small value, it should prevent occurrence of negative pressure calculations. The default value is 1.

36. `gridCoordinates`

The default `gridCoordinates` are `cartesian`, but the `axisymmetric` option is available for 2D translated grids that are to be run in the axisymmetric mode. This option replaces

the pie slice or rotated grid approach for axisymmetric solutions. To use the **axisymmetric** option, the grid must be translated by an arbitrary distance in the z direction, and the solution must be computed in the $x-y$ plane. The mesh cell areas and volumes are adjusted to match the effect of rotating the 2D grid 360° about the x axis, and the translation distance is ignored. The boundary areas reported in the output are for the fully rotated 3D grid when running with **gridCoordinates: axisymmetric**. The correspondence between velocity components in cartesian and cylindrical coordinates is $(u, v, w) \rightarrow (u_z, u_r, u_\theta)$, where $w \rightarrow u_\theta$ is the swirl velocity.

37. coriolis

If this variable is set, then the Coriolis terms are added to the simulation to account for a rotating frame of reference. The rotation is set through the options **axis** that specifies the rotation axis, **center** that specifies the center of rotation, and **speed** that specifies the rotation rate. For example to simulate a moving reference frame that is rotating at 1000 rpm about the x -axis, one would specify:

```
coriolis: <axis=[1,0,0], center=[0,0,0], speed=1000 rpm>
```

38. componentMotion

This variable is used to describe the independent motion of each component in the volume grid. Note: to use this, one needs to have created a grid that contains volume tags using the **vogmerge** utility. The variable contains a list of volume tags assigned to actions, where the actions may be **stationary**, **rotation**, or **prescribed**. For example, to describe a plane with a rotating propeller where the plane is undergoing prescribed motion in a stationary background mesh, one might use the following prescription:

```
componentMotion: < prop = rotation(axis=[1,0,0],center=[0,0,0],speed=400rpm),
                  fuselage = prescribed,
                  background = stationary >
```

The prescribed file will be read from the file **motion_fuselage.dat**, or in general from a file that is of the form **motion_volumeTag.dat**. This file first contains the number of interpolant inputs as the first line, then following for each input there is a list of eight numbers in the following order: *time* (in seconds), *position* (in meters; 3 components, x, y, z), *quaternion* (four components of a normalized quaternion describing rotations about the preceding position). A cubic spline is used to interpolate between the given values. If you want to create a discontinuity in the spline, then repeat the condition for the same time twice.

39. componentHierarchy

This variable is used to establish parent/child relationships between components. For the previous example, this would be used to establish that the **prop** rotation is described relative to the prescribed motion of the fuselage. Therefore, we need to establish that the fuselage is the parent component to the prop. This is accomplished through the following componentHierarchy description:

```
componentHierarchy: < prop = parent(fuselage) >
```

40. componentGeometry

This variable is used to specify the geometry of the hole cut by the enclosing mesh of a component. This is specifically used when performing relative motion simulations where

one component cuts out cells from another components mesh. At least one component will provide the master mesh and will have no component geometry specified. For example, the format for specifying the `componentGeometry` for three bodies in a fixed background mesh is:

```
componentGeometry: < body1 = cylinder(p1=[-1,0,0],p2=[1,0,0],radius=3),
                      body2 = sphere(center=[2,2,0],radius=1.5),
                      body3 = revolution(p1=[5,1,1],p2=[7,1,1],
                      radius=[2.5,4,3],offsets=[0,0.62,1.0]) >
```

Component geometries can be spheres, cylinders, bodies of revolution, or a list of planes. The description of these geometries is as follows:

- **cylinder** Arguments `p1` and `p2` define the two end points of the cylinder axis, while `radius` specifies the cylinder radius. For example, a cylinder with an x axis orientation from $x = 1$ to $x = 2$ with a radius of .5 is given by:

```
cylinder(p1=[1,0,0],p2=[2,0,0],radius=.5)
```

- **sphere** Arguments `center` and `radius` define the sphere's center and radius, respectively. For example, a sphere centered at the origin with a radius of 1.5 is given by:

```
sphere(center=[0,0,0],radius=1.5)
```

- **revolution** This defines a body of revolution about an axis defined by points `p1` and `p2`. The argument `radius` specifies a list of radii at different points along the axis between `p1` and `p2`, while `offsets` provides a list of weighted values for the relative distance between `p1` and `p2` that the associated radius corresponds to.

```
revolution(p1=[50,10,10],p2=[70,10,10],
          radius=[2.5,4,3],offsets=[0,0.62,1.0])
```

- **planeList** This defines a geometry through list of planes. The geometry is defined by the region that is on the left side of all of the planes in the list. Each plane is defined by a point on the plane and a normal vector. The normal vector should point out of the described volume. This method can be used to describe any convex polyhedral region. An example of a cube defined using the `planeList` geometry is given below:

```
// Make a box geometry by defining 6 planes
// this geometry is a cube with edge lengths 2.8m centered on the origin
planeList(list = [
// x coordinate planes
// p is a point on the plane and n is an outward pointing normal vector
plane(p=[1.4,0,0],n=[1,0,0]),plane(p=[-1.4,0,0],n=[-1,0,0]),
// y-coordinate planes
plane(p=[0,1.4,0],n=[0,1,0]),plane(p=[0,-1.4,0],n=[0,-1,0]),
// z coordinate planes
plane(p=[0,0,1.4],n=[0,0,1]),plane(p=[0,0,-1.4],n=[0,0,-1])
])
```

41. componentMRF

The `componentMRF` variable can be used to specify components that are solved in a moving reference frame. This is similar to the `coriolis` term except the velocities are solved in the global reference frame. This can allow simulations with a *frozen rotor* simplification. Note,

this option is incompatible with `componentMotion`. With this variable components can be specified to be either in a rotating or stationary frame of reference. The input syntax is the same as for `componentMotion`.

3.2 Boundary Condition Specifications

The boundary conditions may be assigned as: `inflow`, `supersonicInflow`, `isentropicInflow`, `fixedMass`, `fixedMassOutflow`, `farfield`, `outflow`, `supersonicOutflow`, `extrapolate`, `symmetry`, `impermeable`, `reflecting`, `viscousWall`, `wallLaw`, `periodic`, and `interface`. When a boundary condition requires a specified thermodynamic state, the state specification follows the same format as the `initialConditions` (see Table 1 above).

Notes on boundary conditions:

`inflow`

This is a non-reflecting, characteristic-based boundary condition that uses the given thermodynamic state (*e.g.*, `inflow(p=1atm, T=300K, M=0.3)`) to compute the inflow conditions and the boundary flux. This boundary condition may be used for subsonic or supersonic inflow conditions.

`supersonicInflow`

This boundary condition holds the inflow values at the given conditions but computes a characteristic-based flux at the boundary. This condition can be used for imposing subsonic or supersonic inflow conditions in which the boundary flux is computed from a Riemann solver.

`isentropicInflow`

This inflow boundary condition maintains constant entropy. The user provides the total pressure and total temperature for the associated “plenum” conditions, and the inflow values that preserve the total properties are computed. A typical input is `isentropicInflow(p0=1atm, T0=300K)`. **Note: this boundary condition is only applicable to subsonic inflow conditions.**

`fixedMass`

This inflow boundary condition specifies the mass flow or flux and total temperature. The user provides the `mdot` (kg/s) or `massFlux` ($kg/m^2 - s$) and a total temperature. A typical input is `fixedMass(mdot=0.02 kg/sec, T0=1000 K)`. This boundary condition uses the specified 3D mass rate value for axisymmetric flows in which a 2D grid is translated in the z direction (see the `gridCoordinates` input option). **Note: this boundary condition is only applicable to subsonic inflow conditions.**

`fixedMassOutflow`

This boundary condition specifies the mass flow or flux at an outflow boundary. The user provides the `mdot` (kg/s) or `massFlux` ($kg/m^2 - s$) value only. A typical input is `fixedMassOutflow(mdot=0.02 kg/sec)`. This boundary condition uses the specified 3D mass rate value for axisymmetric flows in which a 2D grid is translated in the z direction (see the `gridCoordinates` input option). **Note: this boundary condition 1) limits the outflow conditions to physically allowable values if the flow becomes choked and 2) extrapolates to the boundary if the flow becomes supersonic.**

farfield

The farfield boundary condition is an inflow/outflow characteristic based boundary condition. It is most suitable for an external flow with a farfield boundary, but it can also be used as a non-reflecting outflow boundary in some internal flows (*e.g.*, `farfield(p=1atm, T=300K, M=0.3)`). It handles both supersonic and subsonic flow conditions.

outflow

This is a characteristic based (subsonic) condition that imposes a constant static pressure at the outflow boundary (*e.g.*, `outflow(p=1atm)`). Use this boundary condition when you intend to hold an outflow boundary at a given pressure. If the outflow becomes supersonic, this boundary condition switches to a supersonic outflow boundary condition (*i.e.*, `extrapolate`) in which the pressure is no longer prescribed. This boundary condition can also be used to establish an average pressure on the outflow boundary if the `pMean` variable is specified. For example, `outflow(pMean=1atm)` would establish an area averaged (mean) pressure on the boundary of one atmosphere.

supersonicOutflow

This boundary condition extrapolates all values to the boundary and should only be used for fully supersonic outflow. If the flow becomes subsonic, it applies a pressure reduction at the boundary to encourage supersonic outflow. This condition does not require any arguments. **Note: This boundary condition should not be used when the outflow may be subsonic.**

This boundary condition extrapolates all values to the boundary. It may be used for supersonic outflow conditions or to establish zero derivatives of velocity, pressure, and temperature in subsonic outflow conditions. This boundary condition switches to `impermeable` if the boundary becomes inflow. This condition does not require any arguments.

symmetry

This boundary condition is specifically designed for 2D or axisymmetric simulations. In these cases, the 2D grid is extruded or rotated to form a one cell thick mesh with two opposing symmetry planes, and this boundary condition is the most efficient and robust. For 3D symmetry planes in which more than one cell separates the planes, `reflecting` is more accurate. Although `symmetry` can be used in these 3D cases, it is only first order accurate. This condition does not require any arguments.

impermeable

This boundary condition implements a second order characteristic-based slip wall boundary condition. Use this boundary for inviscid solid walls on structured meshes. For unstructured meshes, `reflecting` is better behaved. If you must use `impermeable` for unstructured meshes, use it in first order mode by selecting `impermeable(firstOrder)`. The first order mode is also a robust option for `impermeable` on any type of symmetry plane, including those on unstructured grids and 2D/axisymmetric grids with one cell between the planes. This condition does not require any arguments (`firstOrder` is optional).

reflecting

This boundary condition implements a second order reflecting boundary condition where the interior solution is reflected through the boundary normal to the exterior. This boundary condition can be used to implement symmetry planes or impermeable boundary conditions, and it does not require any arguments. **Note: reflecting should not be used to implement symmetry planes when there is only one cell between the planes as it can become unstable.**

viscousWall

This boundary condition implements the no-slip boundary condition for modeling viscous boundary layers. Generally this boundary condition requires a viscous stretched mesh. The **viscousWall** boundary condition may be adiabatic **viscousWall(adiabatic)**, specified heat flux to the fluid **viscousWall(qwall=-1000 watt/m/m)** (the minus sign denotes that the heat flux is into the domain since the boundary normal is positive pointing outward), specified temperature, **viscousWall(Twall=400 K)**, prescribed (see notes on prescribed boundary specifications), or used to interface to other models, **viscousWall(interface)**, for example in conjugate heat transfer. For turbulent computations with a rough wall the sand grain roughness height (k_s in meters) may also be specified, **viscousWall(Twall=400 K, ks = 1e-5)**. **Note: This boundary condition requires normal grid spacing of $y^+ \sim 1$ to achieve accurate solutions.**

periodic

This boundary condition is used to specify a periodic condition between pairs of boundaries which are identified using the same name. One periodic boundary specification in the pair specifies the transformation between the boundaries, which may be either translation or rotation. Translation is specified through a translation vector, and rotation is specified by a center of rotation, an axis of rotation, and a rotation angle.

```
// Translated periodic boundary condition
BC_1=periodic(name="A", translate=[.8,0,0]),
BC_2=periodic(name="A"),
// rotated periodic boundary condition
BC_3=periodic(name="B", center=0, vector=[1,0,0],
            rotate=30 deg),
BC_4=periodic(name="B"),
```

turboInterface

The **turboInterface** is a specialized paired interface that was developed specifically for turbomachinery simulations. In the turbo interface boundary condition surfaces are paired by assigning two different boundaries with the same name. For example, to connect a stator section to a rotor section one may couple these two boundary conditions using **turboInterface** using the specification:

```
statorOutflow = turboInterface(name="interface1"),
rotorInflow = turboInterface(name="interface1"),
```

The two boundary conditions are paired through sharing the name “interface1”. There are additional controls that can be provided in the vars file to control the behavior of the **turboInterface** boundary condition. Note that the correct function of this boundary condition requires that exactly two boundaries share the same name. Multiple boundaries can be paired and identified with independent names. The names given to each pair are arbitrary and can be selected by the user based on the problem at hand. They are listed as follows:

- **turbo_walk_history_start**

This variable gives how many steps a particle walks before it will start remembering its path to prevent cycles in the geometric search. The default value is 9.

- `turbo_max_walking_warning`

This variable gives the number of steps that a particle will walk before issuing a warning and terminating its search. The default setting is 500.

- `turboSearchGeometry`

This variable gives the type of search used to find cells for interpolation. it can be set to either 'cartesian' or 'cylindrical'. If it is set to cylindrical, then it is assumed that the axis of the cylindrical coordinate system is the x axis. The default is to use cartesian searching. For turbomachinery applications cylindrical searching can provide more accurate interpolations for boundary layer regions near viscous walls.

3.2.1 Notes on Vector Value Inputs

Vector values (such as inflow boundary condition velocities or directed Mach numbers) can be input in both Cartesian and spherical coordinates. By default, if only a single value is assigned to the vector, it is assumed that this value points in the x coordinate direction. A Cartesian vector is assigned by providing a list of three values for the x , y , and z components, respectively, enclosed in square brackets, `[]`. For example:

```
(u = [0,20 m/s, 0]) // Velocity in the +y direction
```

The magnitude of the vector and the keyword `normal` may also be used to specify that the velocity is normal to the boundary:

```
(u = 20 m/s,normal) // Velocity normal to the boundary
```

Likewise, we can specify the velocity using a polar notation of `polar(mag, Θ , Ψ)` where the final vector is given by $u = [mag * \cos(\Theta) * \cos(\Psi), mag * \sin(\Theta) * \cos(\Psi), mag * \sin(\Theta) * \sin(\Psi)]$. For example:

```
(u = polar(20 meters/second,5 deg, 0)) // 5 degrees angle of attack
```

The definition of the polar angles and the resulting u , v , and w velocity components are shown for a right-handed Cartesian coordinate system in Fig. 1.




Figure 1: Definition of the polar angles used to specify the velocity components.

3.2.2 Notes on Inflow Velocity Specifications

A typical inflow boundary might be specified as:

```
inflow(p=5000.0 psi,T=273 K,u=100, normal)
```

but there are also two possible specifications for rotating flows: rigid body rotation (helical flows) and a fixed swirling flow angle. Rigid body rotation is specified using the `rotAxis`, `rotCenter`, and `rotSpeed` boundary condition options, where `rotAxis` is a vector that defines the axis of rotation, `rotCenter` is a point on the rotation axis, and `rotSpeed` is the rotational speed given by default in revolutions per minute (rpm). The rotation is defined by the right-hand rule, thus a positive rotation speed about the x -axis `rotAxis` will produce positive z components of velocity when the value of y is positive.

Alternatively, it is possible to specify a fixed swirl angle whereby the direction of the velocity will rotate about a center axis. This is done using the specification of `swirlAxis`, `swirlCenter`, and `swirlAngle`. The specification is similar to the rotating flow specification, with the specified rotating flow conforming to a right-handed coordinate system.

Both rigid body rotation and swirling flows can be specified for any inflow boundary condition (`inflow`, `supersonicInflow`, `farfield`, `fixedMass`, and `isentropicInflow`). When these conditions are specified, the direction of the velocity vector is modified to conform to the rotating conditions, but the magnitude of the velocity is preserved. For `fixedMass` and `isentropicInflow`, best results will be achieved if the boundary surface is planar and normal to the axis of rotation. The following are some examples:

```
inflow(p=5000.0 psi, T=274 K, u=100 m/s,
      swirlAxis=[1,0,0], swirlCenter=[-1,0,0], swirlAngle=5deg),
fixedMass(mdot=1kg/s, T0=300K,
          rotAxis=[0,1,0], rotCenter=[0,0,0],
          rotSpeed=3000 rotations/minute),
```

3.2.3 Notes on Setting Viscous Boundary Velocity

A tangential velocity can be specified on viscous wall boundaries for simulating wall translation and rotation. To specify the translational velocity, use the `Uwall` parameter. A rotating viscous wall is specified using the `rotAxis`, `rotCenter`, and `rotSpeed` options, similar to rotating inflows. A moving wall velocity can be specified for either `viscousWall` or `wallLaw` boundary conditions. In addition, only the tangential part of the specified velocity is used in the viscous velocity specification.

Examples

A moving lid with a constant linear velocity is specified as:

```
wallLaw(adiabatic, Uwall=[1,0,0]), // Wall moving at 1 m/s to right
```

A rotating wall can be specified as:

```
viscousWall(Twall=300K, rotSpeed=100 rpm,
            rotAxis=[0,0,1], rotCenter=0),
```

where the rotation axis is the z axis.

3.2.4 Integrated Forces and Moments

FlowPsi provides the integrated forces and moments about a moment reference center specified by the user (the default moment center is the origin of the coordinate axes). The output for specific BCs as well as all of the viscous walls is placed in the standard output as well as in special files located in the `/output` directory. The output file naming convention and the file format are discussed below.

The user specifies a moment reference center in the `.vars` file for each BC over which integrated forces and moments are required. For example,

```
BC_4 = viscousWall(adiabatic, momentCenter = [1., 0.5, 0.])
```

locates the moment reference center for BC_4 at (1.,0.5,0.) meters (the default units are meters).

Both the inviscid pressure and viscous shear stresses are integrated over each BC. For axisymmetric flows in which a 2D grid is translated in the z direction¹, the integration is over the 3D surface area resulting from a 360° rotation about the x -axis, and the integrations over the boundary area are computed accordingly. For the BCs on which a `momentCenter` is specified, *FlowPsi* also computes the moments for each of these force contributions. The forces and moments are resolved into the three coordinate directions (x, y, z). Note that for the inviscid pressure contribution, the force actually integrated is $(p - p_{ref})$, where p_{ref} is the user input gauge pressure (`.vars` file input `p0`) (default $p_{ref} = 0$).

The integrated fluxes for a specific boundary are placed in a file named `/output/flux_BC4.dat`, for example. The output format in this file is:

```
iteration, simulation time, mass flux, Fx, Fy, Fz, energy flux, area
```

where `Fx`, `Fy`, `Fz` are the total momentum fluxes (viscous plus inviscid) in each coordinate direction. Similarly, the moments are placed in a file named `/output/moment_BC4.dat`, and the file format is:

```
iteration, simulation time, Mx, My, Mz, RefMx, RefMy, RefMz
```

where `Mx`, `My`, `Mz` are the actual moments ($N\cdot m$) relative to the defined `momentCenter`, and `RefMx`, `RefMy`, `RefMz` are the reference moments based on the integration of $(p - p_{ref})$ about the `momentCenter`.

The force integrations over each BC surface are defined as follows:

$$F_x = \int d\mathbf{F} \cdot n_x dA \quad (1)$$

$$F_y = \int d\mathbf{F} \cdot n_y dA \quad (2)$$

$$F_z = \int d\mathbf{F} \cdot n_z dA, \quad (3)$$

where $d\mathbf{F}$ is the sum of the viscous and inviscid momentum fluxes on a surface with area dA , and (n_x, n_y, n_z) are the components of the surface normal vector. The moment integrations are:

$$\mathbf{M} = - \int (\mathbf{r} - \mathbf{r}_0) \times \mathbf{F} dA, \quad (4)$$

¹see the `gridCoordinates` input option for further information on computing axisymmetric flows with 2D translated grids.

where \mathbf{r}_0 is the user specified `momentCenter` vector, and \mathbf{r} is the face center location of dA . The *flowPsi* output contains the (x, y, z) components of the total moment vector \mathbf{M} . For axisymmetric flows in which the 2D grid is translated in the z direction, F_x and F_y take on the integrated value while $F_z = 0$. Similarly, $M_x = M_y = 0$, and M_z takes the integrated value. Note that the moment integration is consistent with the aerodynamic convention for pitch, roll, and yaw (*i.e.*, positive moments occur for pitch up, roll to the right, and yaw to the right when viewed from the pilot's perspective).

4 Relative Motion Problems using *FlowPsi*

FlowPsi can be used to simulate moving bodies using an unstructured over-set technique. The over-set technique can be used with or without hole-cutting. Usually for relative motion in turbo-machinery applications, hole cutting can be avoided through careful design of the mesh components. The most important tool in developing simulations with relative motion is the interface boundary condition. This boundary condition is applied at the boundary faces that are at the edge of two grid regions. Values are interpolated automatically from neighboring meshes at the interface boundary condition. Generally this boundary condition can be used for either overlapping meshes or meshes that just slide along a common surface (but are not point matched).

In order to move components independently, one must first create a grid that contains volume tags that identify the meshes associated with each component. Each component mesh is generated separately, and then all of the meshes are merged using the `vogmerge` utility. During the merging step, one provides volume tags associated with each component grid. For example, to merge the rotor and stator components into a single grid from two grids called 'rotor.vog' and 'stator.vog', one would execute the following command:

```
vogmerge -g rotor.vog -tag Rotor -g stator.vog -tag Stator -o pump.vog
```

This command combines the two grids for the rotor and stator into a single grid called 'pump.vog'. The cells of the pump grid that come from `rotor.vog` are tagged with "Rotor", and the cells of the pump grid that come from `stator.vog` are tagged with "Stator".

Once a grid with the appropriate embedded volume tags and interface boundary surfaces are defined, it is possible to move each component independently using the `componentMotion` variable in the `vars` file. If this variable is defined, then the motion of each component will be specified by the user (note that if no `componentMotion` variable exists, then one can still use the overset boundary conditions, but all of the mesh components will remain stationary.). There are several options for component motion:

1. **stationary**

Use this option to indicate that a given component does not move in the simulation.

2. **rotation**

This option is used to define the fixed rotation rate of the component. The arguments to this option are **axis**, which specifies the axis of rotation, **center**, which specifies a point on the rotation axis, and **speed**, the rotation rate in revolutions per minute (rpm).

3. **prescribed**

This option is used to specify that the component motion is given by a user defined data file. This selection has no arguments. The file name is derived from the component tag name, and the naming convention is `motion.TAG.vog`, where **TAG** is the tag name given to the component in the `vogmerge` step. The file starts with an integer number that specifies

the number of data entries in the file. The data entries in the prescribed motion file (one entry per line) then consist of time in seconds, followed by the three coordinate positions of the center of mass in meters, followed by the four components of a quaternion that describes rotations about the center of mass. Interpolation between the given data is accomplished using a cubic spline. In order to specify derivative discontinuities in the data, repeat a data entry twice.

4.1 Specifying motion in a relative frame

In some cases, one will specify the motion of a component relative to another component's reference frame. For example, one may describe the rotation of the propellers of an airplane relative to the airplane's coordinates. If the airplane follows a prescribed motion, it is naturally expected that the propellers will follow the same prescribed motion. In order to specify that the above motion is described in a relative frame, one uses `componentHierarchy` to define the parent child relationship between objects. If there is no motion prescribed in a frame relative to another, then this option is not needed. Details on the usage of `componentHierarchy` are given in Section 3.1.

4.2 Hole Cutting

When motion can be prescribed such that it occurs along sliding, non-matched surfaces or along overlapping interface boundary conditions that do not intersect other boundary faces, then hole cutting can be avoided. This reduces the number of cells needed and improves run times. However, it is impossible to describe more general motion, such as two bodies that come in close proximity, without hole cutting. To employ hole cutting, a grid is first created around an object that will be moving in the domain. Usually the outer envelope of this grid (where the interface BC is applied) will be a simple geometric shape such as a sphere, cylinder, or body of revolution. Hole cutting is enabled by adding a definition for `componentGeometry` to the `.vars` file. (Note: adding `componentGeometry` to the `.vars` file enables the hole cutting feature.). In the component geometry input, the geometry of the outer envelope of given mesh components will be specified. At least one of the mesh components will not be given any geometry description as this mesh will serve as the background mesh.

When employing hole cutting, one does not need to be concerned with the overlap of the outer interface boundary of a component since the hole cutting will automatically cut cells according to the distance of the cell to surfaces from other tagged regions. Thus when the interface surface overlaps geometry from another grid, the grids will be cut along the surface of equal distance between the two component surfaces.

5 FlowPsi Pre- and Post-processing Utilities

5.1 FlowPsi Pre-processing Utilities

5.1.1 ugrid2vog

This converter allows one to convert *SolidMesh* or *aflr3* mesh files to the volume grid representation used by *FlowPsi*. This command can be run in parallel using *mpirun* for faster processing of large meshes. By default, the command will detect if the input file is binary or ASCII, but you can force it to read a binary `.b8.ugrid` file using the `-b` option. The converter also has a `-o` option to disable optimizations that would change the numbering of nodes and cells in the mesh from the given `ugrid` file. Generally this option would only be used if it is important for the grid conversion to preserve the original node ordering. Additionally, you may input the units of the

ugrid mesh (default is meters) with the option `-in` for inches, `-ft` for feet, `-cm` for centimeters, or `-mm` for millimeters. If another unit is needed, you may use the `-Lref` option. For example, if the input grid named "wingsection.b8.ugrid" is non-dimensionalized with respect to a reference length of 1.57 inches, you would use the command:

```
ugrid2vog -Lref "1.57 inch" wingsection
```

5.1.2 xdr2vog

This converter allows one to convert the older XDR format mesh files to the volume grid representation used by *FlowPsi*. This command can be run in parallel using `mpirun` for faster processing of large meshes. The converter has a `-o` option to disable optimizations that would change the numbering of nodes and cells in the mesh from the given XDR mesh file. Additionally, you may input the units of the XDR mesh with the option `-in` for inches, `-ft` for feet, `-cm` for centimeters, or `-mm` for millimeters. With no option, `xdr2vog` will assume the input is in meters. The `-Lref` option can be used as with `ugrid2vog` to specify other unit systems.

5.1.3 vogmerge

The `vogmerge` utility allows one to combine VOG file grids, apply geometric transformations to the grids, rename boundary tags, and apply volume tags. Use options

```
vogmerge -g input.vog -zscale -1 -g input.vog -o combine.vog
```

Note that the geometric transform options are applied to the grid of the preceding `-g` options. If many geometric transforms are specified, then they are applied in the given order. Finally, for each grid, one may apply a volume tag using the `-tag` option. This is used when creating grids for simulations using relative motion. For example, to perform a simulation on turbo-machinery that included a rotor and stator section, one would generate two grids, one for each component, and then merge them using the following `vogmerge` command:

```
vogmerge -g rotor.vog -tag rotor -g stator.vog -tag stator -o turbo.vog
```

Once can then prescribe different motions to each grid using the `componentMotion` variable in the `.vars` file using the volume tags `rotor` and `stator`.

Executing `vogmerge` without any options will provide a usage guide. The current options to `vogmerge` are:

5.1.4 cobalt2vog

The `cobalt2vog` grid converter will convert an ASCII unformatted *Cobalt 60* grid format file and convert it into a VOG grid suitable for use with *FlowPsi*. The extension on the cobalt grid should be `cog`.

5.1.5 plot3d2vog

The `plot3d2vog` grid converter will convert multi-block `plot3d` files into an VOG file format compatible with *FlowPsi*.

5.1.6 fluent2vog

The `fluent2vog` grid converter will convert fluent `.msh` files to the VOG file format for use with *FlowPsi*.

option	description
-g <i>gridfile</i>	Specify input grid filename
-o <i>gridfile</i>	Specify output grid filename
-xshift <i>value</i>	translate grid in x-coordinate
-yshift <i>value</i>	translate grid in y-coordinate
-zshift <i>value</i>	translate grid in z-coordinate
-xscale <i>value</i>	scale grid in x-coordinate
-yscale <i>value</i>	scale grid in y-coordinate
-zscale <i>value</i>	scale grid in z-coordinate
-scale <i>value</i>	scale grid in all coordinates
-xrotate <i>value</i>	rotate about x-axis (degrees)
-yrotate <i>value</i>	rotate about y-axis (degrees)
-zrotate <i>value</i>	rotate about z-axis (degrees)
-bc <i>oldname,newname</i>	rename boundary surface
-glue <i>bname</i>	glue this boundary to a matching bc
-tag <i>name</i>	specify volume tag for input grid

5.1.7 ccm2vog

The *ccm2vog* grid converter will convert Star CCM+ grids to the VOG file format for use with *FlowPsi*.

5.1.8 vogcheck

The *vogcheck* utility is a grid quality assessment tool for *FlowPsi*. This program can be run in parallel for larger grids using *mpirun*. The program outputs an overall assessment of the grid quality in the range of **UNUSABLE**, **marginal**, **poor**, **good**, and **excellent**. A grid that is identified as **UNUSABLE** will probably not be useful for *FlowPsi* simulations. Grids marked as **marginal** may provide numerical robustness problems for *FlowPsi*. It is suggested that the user query the cause of poor mesh quality for these grids. Grids identified as **poor** may be acceptable for running with *FlowPsi*, but accuracy will probably be degraded in some regions. Any grid marked as **good** or **excellent** has a mesh quality that is suitable for high quality solutions using *FlowPsi*.

To query the mesh quality, one can extract the mesh quality measures from *vogcheck* from iteration 0.

5.2 FlowPsi Post-processing utilities

5.2.1 Probe Output

The *FlowPsi* code can output variables at a specified location to a file every iteration using probes. There is no limit to the number of probes allowed. The specification in the *.vars* file is:

```
probe: <probe1=[0,0,0],probe2=[1,0,1]>
```

where the units of the point are in meters if units are not explicitly specified with the point locations. Probes will generate output files in the form **probeX.dat**, where **X** is the probe number. The file contains lines with the following information:

```
<ncycle> <stime> <T> <p> <rho> <a> [ <U> ] [ <pos> ] <dist> <mass fractions>
```

where:

`<ncycle>` is the iteration number of the simulation
`<stime>` is the simulation time
`<T>` is the temperature
`<p>` is the pressure
`<rho>` is the density
`<a>` is the sound speed
`<U>` is the velocity vector
`<pos>` is the position of the probe
`<dist>` is the distance between the probed value and the position given in the input file.

5.2.2 Plot Format Converter

See the `plot_freq` item for an explanation of how to generate data for visualization in the `output` directory. These files are annotated by the iteration number (modulo the `plot_modulo` value). Once generated, the `extract` program can be used to generate various plot files for post processing programs such as *FieldView*, *EnSight*, *TecPlot*, and *2dgv*. The first argument to `extract` specifies which post processing software format to use. Table 2 gives a list of the currently supported post processing selectors for `extract`. Other options may follow that are specific to the particular post processor. The user then provides the problem name (this is the name of the *FlowPsi* `.vars` file without the `.vars` extension), the iteration number to extract from, and a list of variables to extract. The default variables output by *FlowPsi* are given in table 3. The *2dgv* post processing package will only allow the specification of a single variable, while all others support visualization of multiple scalar and vector values. The form of an `extract` command is:

```
extract -<package> [options] <problem name> <time step> <variable(s)>
```

Note that documentation of how to use `extract` can also be obtained by executing the command `extract` without any arguments.

post-processor option	purpose
<code>-fv</code>	<i>FieldView</i> postprocessing
<code>-en</code>	<i>EnSight</i> postprocessing
<code>-tec</code>	<i>TecPlot</i> postprocessing
<code>-2d</code>	<i>2dgv</i> postprocessing
<code>-cut</code>	Dump a cutting plane for <i>2dgv</i> postprocessing
<code>-ascii</code>	Dump selected nodal data to an ASCII file
<code>-mean</code>	Generate mean and variance from family of variables
<code>-surf</code>	Extract boundary surface mesh

Table 2: `extract` post-processor selectors

The *FlowPsi* code will optionally output other scalar and vector data when these variables are provided in the `plot_output` list in the `.vars` file. For example, to extract the turbulence model scalars `k`, `w`, and `tmu`, one would place the following in the `.vars` file:

```
// Note, no spaces between variables, just commas
plot_output: k,w,tmu
```

Now the variables can be extracted for visualization with *FieldView* for the case `plume` at iteration 100 using the command:

```
extract -fv plume 100 k w tmu
```

Key	function
r	density
p	$\log_{10}(\text{pressure})$
P	actual pressure in SI units
u	fluid speed
0, 1, 2	Cartesian components of fluid velocity
v	velocity vector
m	Mach number
t	temperature
a	sound speed
tau	wall shear stress (vector)
taumag	wall shear stress magnitude (scalar)
qdot	wall heat flux
tw	wall temperature
pw	wall pressure
bt	boundary temperature
bpg	boundary gage pressure
br	boundary density
mdot	boundary mass flux

Table 3: **extract** command keys

5.3 Surface Data Output Options

By default *FlowPsi* outputs data for volume nodal data and surface facet information. For very large meshes or time accurate simulations where frequent output is required, the size of the volume data can become prohibitive. To facilitate more control over the amount of output data *FlowPsi* produces, users can instead choose to write out surface data for a specific subset of surfaces, variables, and times. By default individual surface data is not written out, but the output for any boundary surface can be enabled by adding the variable **plotFreq** and **plotVars**. The variable **plotFreq** tells the frequency of plotting for that boundary surfaces whereas **plotVars** is set to a list of variables that will be output for that boundary surface. For example, if one wanted to output surface pressures and temperatures every 10 cycles for a viscous wall boundary one could enter in the **boundary_conditions** line:

```
bladerow = viscousWall(adiabatic, plotFreq=10, plotVars=[pg,t]),
```

In addition to boundary surface data, *FlowPsi* can also output clip surfaces in the form of clip planes or isosurfaces. The clip surface capability is enabled by adding the variables **clipFreq** and **clipSurfaces** to the vars file. The **clipFreq** variable provides an integer which will be the frequency of clip surface output, while the **clipSurfaces** gives a list of surface names with clipping specifications. The clipping specification can either be a **plane** or an **isosurface**. For a clip plane the specification requires a **point** on the plane and the **normal** vector of the plane. The iso-surfaces specifications gives a variable assignment where the iso-surface will be computed. An example of clip surfaces is given in the specification:

```
clipFreq: 25
clipSurfaces: < zcut10 = plane(point=[0,0,10],normal=[0,0,1]),
               t1000 = isosurface(t=1000),
               t0500 = isosurface(t=500)>
```

The previous example will generate cut surfaces every 25 timesteps where the cut surfaces include the surface named “zcut10” which is the plane formed by $z = 10$, the “t1000” isosurface where temperature is 1000 Kelvin and the “t500” isosurface where temperature is 500 Kelvin. Note, the input for the isosurface command cannot accept unit specification and instead will always be interpreted in terms of *FlowPsi*’s output default which is the MKS unit system. Also note, any variable that was specified for plotting on the volume mesh (those listed in `plot_output` or `plot_output_exclusive` will be written out for these surfaces.

5.4 Mean and Variance Output

For hybrid RANS/LES or DES simulations, it is useful to compute mean and variances of output variables. There are two methods provided to allow users to extract this type of information from a simulation: offline and online. The offline approach is implemented as part of the `extract` utility with the `-mean` option. The `extract` utility can read in a sequence of output scalars or vectors that were written out by *FlowPsi* and compute mean and variance values which are then written back out to the output directory. The `extract` utility can then be used to visualize the resulting mean and variance values. To extract these values, use the argument `-mean` followed by `-end` with the ending iteration and `-inc` with the iteration increment. `Extract` starts with the usual given `extract` iteration number and averages consecutive plot files at the provided increment until the ending iteration is reached. The mean values are then written to the output directory as the variable with `Mean` appended for the mean and `Var` for the variance. For vector valued variables, the covariance is also output as `Cuv`, `Cuw`, and `Cvw` for the covariance of the x - y , x - z , and y - z components, respectively. These variables can then be extracted in the usual way for visualization.

The online method for mean and variances is enabled by adding the following line to the top of your `.vars` file (before the first brace, `{`):

```
loadModule: onlineStats
```

In the `.vars` file, you will also need to add a definition for the averaging window size using the `meanFreq` variable. The average will automatically reset when the window size is reached. Alternatively, the code can perform exponential averaging if the variable `useExponentialMean` is set to `on`. In the exponential averaging approach, the averaging weight is computed based on the window provided by `meanFreq`. When online averaging is enabled, the code will output the variables listed in table 4. Note that these variance values are the square of the RMS value. Also note that the velocity mean and variances are Favre averaged while

Key	function
<code>tmean</code>	average temperature
<code>tvar</code>	variance in temperature
<code>pgmean</code>	mean gage pressure
<code>pgvar</code>	variance in gage pressure
<code>rmean</code>	mean density
<code>umean</code>	Favre mean velocity vector
<code>uvar</code>	Favre variance in velocity components
<code>ucovar</code>	Favre covariance of velocity components

Table 4: variables output by the `onlineStats` module

5.5 The Grid Motion Module

The grid motion module smoothly projects deformations from surfaces into the volume mesh. The algorithm used by the grid motion module employs an algebraic interpolation method that is fast and robust. High speed is obtained through a tree code acceleration which allows the interpolation to be evaluated to within a specified accuracy.

5.5.1 Theoretical Background

The deformation of the volume mesh can be viewed as a projection of deformations and rotations from the surface into the volume. For general mesh deformations the local surface rotations can be estimated once the displacements are given for each node of the surface mesh. In the algorithm described in this paper we estimate local rotations by finding the closest rotation about the node that best matches the relative displacements of all edges and normals from surface facets that reference the given node. Using this information, the displacement field presented by any surface node can be described as a rigid body motion. In our approach, the global deformation field is given as a weighted sum of these nodal rigid body displacements, where node i on the deforming surface presents a displacement field, denoted $\vec{s}_i(\vec{r})$, that is written as

$$\vec{s}_i(\vec{r}) = M_i \vec{r} + \vec{b}_i - \vec{r}, \quad (5)$$

where M_i is a rotation matrix, \vec{b}_i is a displacement vector associated with the i^{th} node, and \vec{r} is a coordinate vector in the original mesh. The displacement field in the volume mesh is then described through a weighted average of all boundary node displacement fields as given by

$$\vec{s}(\vec{r}) = \frac{\sum w_i(\vec{r}) \vec{s}_i(\vec{r})}{\sum w_i(\vec{r})}. \quad (6)$$

For the interpolation weight function, we choose a weighting that is a function of the reciprocal of distance. We utilize a two-exponent form of the weighting function so that we can preserve near-boundary deformations while providing a smooth transition in the interpolation to blend deformations in the bulk of the volume mesh. We also use the area of the surface facet in the weighting function so that mesh refinement of a region does not increase its influence in the interpolation method. The resulting weighting function is given as

$$w_i(\vec{r}) = A_i * \left[\left(\frac{L_{def}}{\|\vec{r} - \vec{r}_i\|} \right)^a + \left(\frac{\alpha L_{def}}{\|\vec{r} - \vec{r}_i\|} \right)^b \right], \quad (7)$$

where \vec{r}_i is the position of point i , A_i is the area weight assigned to node i , L_{def} is an estimated length of the deformation region, α is an estimated size of the near body influence region expressed as a fraction of L_{def} , and a and b are user-defined exponents. Numerical experimentation suggests that $a = 3$ and $b = 5$ provide the best-quality results. The L_{def} parameter may be specified by the user, but it is usually determined automatically by computing the maximum distance any mesh node is from the mesh centroid. The alpha parameter controls how much weight is given to deformations of nearby surfaces relative to more distant ones. Generally, this will need to be increased with the variability in surface displacements. This appropriate automatic determination of α can be computed from the deviation of the boundary displacement field from the mean boundary displacement given by

$$\vec{s}_{mean} = \sum_{i=1}^N a_i * \vec{s}_i(\vec{r}_i), \quad (8)$$

where $a_i = A_i / \sum_{j=1}^N A_j$ is the normalized node weight and α is then determined through the following expression:

$$\alpha = \frac{\eta}{L_{def}} \max_{i=1}^N \|\vec{s}_i(\vec{r}_i) - \vec{s}_{mean}\|. \quad (9)$$

The η term is a user defined parameter that expresses how far into the domain a displacement should be “absorbed.” Numerical experiments have found that a value of $\eta = 5$ provides satisfactory results in most cases. Generally, we limit α to be 0.1 or greater to guarantee that a region of rigid deformation is maintained in boundary layer regions to preserve viscous mesh quality. This defines our baseline direct interpolation algorithm. Our numerical experiments show that this approach is able to produce high-quality meshes for large deformations and also is able to preserve mesh orthogonality near the boundary surfaces.

5.5.2 Optimizations

For two-dimensional problems where the number of boundary nodes is relatively small, the above algorithm runs quite fast. However, for three-dimensional problems, the surface nodes grow at a rate of $O(n^{\frac{2}{3}})$, where n is the total number of nodes in the simulation. Therefore, the cost of the above algorithm becomes $O(n * n^{\frac{2}{3}}) = O(n^{\frac{5}{3}})$. Although the algorithm parallelizes well, the cost of the deformation calculation quickly becomes dominant for three-dimensional cases. However, for mesh deformation, the accuracy of this evaluation is not critical and only needs to be accurate to a fairly large fraction of the mesh spacing. Also, the form of the sum is similar in form to gravitational potential equations that are solved for N-body simulations. It is therefore reasonable to look towards tree-code optimizations similar to N-body algorithms such as Barnes-Hut[3] and FMM[4]. Using these approaches, the cost of the deformation step could be reduced to $O(n \log n)$, or even $O(n)$ if suitable kernels for the multipole and local expansions are found.

The Barnes-Hut tree-code algorithm requires that we have an approximation for the influence of a set of points at some distance and that we have an estimate for the error in the approximation. In the case of N-body simulations, potential theory can be used to express the influences of a collection of bodies as a multipole expansion about the centroid. A truncated expansion can be used to provide an approximation and also an estimate on error bounds. Unfortunately, the sums given in equations (6) and (7) do not satisfy potential theory and so an analytical approach to forming a multipole expansion is not readily available.

Lacking this infrastructure, we find an alternative way of constructing a suitable approximation to the effects of a collection of particles. First, we focus on approximating the sum

$$f(\vec{r}) = \sum_{i=1}^N \frac{a_i}{\|\vec{r} - \vec{r}_i\|^3}, \quad (10)$$

where a_i is the normalized node weight. Instead of attempting to find a multipole expansion for this sum, we instead focus on finding a small set of pseudo nodes that can approximate the collection of particles at a distance, essentially approximating a truncated multipole expansion. For example, to approximate a quadrupole we would use four points that share the same centroid as the collection of points such that our sum is approximated by

$$f(\vec{r}) \approx \frac{1}{4} \left(\frac{1}{\|\vec{r} - \vec{q}_1\|^3} + \frac{1}{\|\vec{r} - \vec{q}_2\|^3} + \frac{1}{\|\vec{r} - \vec{q}_3\|^3} + \frac{1}{\|\vec{r} - \vec{q}_4\|^3} \right), \quad (11)$$

where $\vec{q}_1, \vec{q}_2, \vec{q}_3$, and \vec{q}_4 represent the locations of the four approximating pseudo nodes which we call the quad-points or QPs. Given the constraint that the QPs share the centroid of the collection given in equation 10, we can determine the location of \vec{q}_4 using the relation

$$\vec{q}_4 = 4 \left(\sum_{i=1}^N a_i \vec{r}_i \right) - (\vec{q}_1 + \vec{q}_2 + \vec{q}_3). \quad (12)$$




Figure 2: Illustration of QP computation

The three remaining QPs are found using a non-linear least squares that minimizes the error given by

$$err = \|f(\vec{r}) - \frac{1}{4} \left(\frac{1}{\|\vec{r} - \vec{q}_1\|^3} + \frac{1}{\|\vec{r} - \vec{q}_2\|^3} + \frac{1}{\|\vec{r} - \vec{q}_3\|^3} + \frac{1}{\|\vec{r} - \vec{q}_4\|^3} \right)\|^2 \quad (13)$$

on the test sphere defined by $r = 3R_e$ where R_e is the radius of the smallest centroid centered sphere that contains all of the points, \vec{r}_i . In our case we minimize the error at the 20 nodes of a dodecahedron that is circumscribed by the test sphere. We pick three points arbitrarily from the collection of points to initialize the non-linear least squares solve. An example of this procedure is shown in figure 2 where the small white points show a collection of 128 random points sampled from the surface of one octant of the unit sphere. The dark gray inner region illustrates the enclosing sphere with radius R_e , while the larger sphere illustrates the test sphere where error is minimized. The larger blue points show the location of the computed QPs. Once we have computed the QP locations, we utilize 12 additional points on the sphere defined by the centroids of the dodecahedron faces to sample the errors on the test sphere. The errors in using the QPs to approximate the sum for both the third and fifth powers are shown in table 5. The table shows that the approximation error drops as a function of $1/r^2$ for exponents three and five providing the needed model for approximation errors.

For general deformations, we also need to consider the implications of the variation of the rotation matrices, M_i , and translation vectors, b_i . To approximate the effects of these terms, we use the QPs found in the previous step, but perform a term-by-term least squares fit in order to assign M_i and b_i to the four approximation points, again preserving the centroidal value. The resulting approximations have similar accuracy profiles to those shown in table 5.

From these approximations, we have the foundation to build a tree-code fast evaluator. The architecture of the tree-code mesh deformation software is shown in figure 3. The boxes with double lines only need to be evaluated once using the original undeformed grids. The remaining steps, because they depend on the actual displacements, must be computed every time new displacements are provided. The first step is to build the KD-tree to spatially partition the surface nodes. To reduce the depth of the tree, we use bins containing approximately 32 surface nodes for the leaves of the tree. Once the spatial partition is established we then compute the

Table 5: Errors for QP approximation

distance (r)	error (a_i/r_i^3)	error (a_i/r_i^5)
$3R_e$	0.967%	3.34%
$10R_e$	0.0563%	0.140%
$20R_e$	0.0124%	0.0313%
$40R_e$	0.00290%	0.00737%
$80R_e$	0.000717%	0.00180%
$160R_e$	0.000184%	0.000457%




Figure 3: Mesh Deformation Software Architecture

QPs for each vertex in the KD-tree. Once the QPs are determined, the errors at the test sphere can be measured. The KD-tree vertices are then populated with the centroid, containing sphere radius, and measured errors.

The final step in the deformation algorithm is evaluating deformations for the mesh volume nodes (note the deforming surface nodes are evaluated using the prescribed deformations). The recursive evaluator computes the numerator, denominator, and the error in the denominator of equation (6). If we are a leaf of the tree then the deformation and weights are evaluated exactly forming the base case of the recursive algorithm. Otherwise we evaluate the error associated with using the QP approximation at the current level. If the error bound is met, then return the approximate numerator and denominator, otherwise sum the evaluation of the two children. Note, since we compute the accumulated error, there can be an advantage to evaluating the parts of the tree that give larger weights as this increases the error threshold for the remaining part of the tree. Therefore we visit child vertices in the order of closest centroid first to reduce computational cost.

The behavior of this algorithm automatically reduces errors near the surface since in this region the nearby leaves will be evaluated as an exact expansion. Since the weights associated with these nearby points will typically be much greater than more distant points, the error in the boundary layer regions is generally very small regardless of target error bound. Further away from the surface we generally only need the deformation error to be within small fraction of the edge length. Typically we find that this requirement is easily met with a target relative error of just 1%.

5.5.3 Grid Motion User Controls

The grid motion module is activated by adding an appropriate `loadModule` command at the top of your vars file:

```
loadModule: gridMotion
```

The global properties of the interpolation can be controlled by user specified parameters. In most cases these parameters will work very well in their default setting, but in specific cases it may be helpful to make changes. These controls are represented by the following vars file variables:

1. `gridMotionExpA`

This is the a coefficient described in the weight equation, eq. (7). This controls how deformations are blended away from the deforming surfaces. The exponent is a positive integer with a default setting of 3. A smaller setting will make the blending smoother with a greater risk of having grid line crossing, whereas a larger number will increase shearing in the mesh. Generally it is good to leave this value as the default.

2. `gridMotionExpB`

This is the b coefficient described in the weight equation, eq. (7). This controls how deformations are blended for points close to the surface. It is usually set to a larger value than a so that the behavior near the body behaves in a more rigid manner preserving viscous mesh orthogonality. The default setting for this is 5.

3. `gridMotionLref`

This is the L_{def} coefficient in the weight equation, eq. (7). This defines the average length over which the deformation will be distributed. If this variable is not set, then the deformation length is computed by finding the most distant point in the mesh from the mesh centroid. For many problems this is a good default value, but for specific cases, particularly

where you know the deformation region is much smaller than the overall mesh, it may be helpful to set this value to something that is sensible to your specific mesh deformation case.

4. `gridMotionAlpha`

This is the α coefficient in the weight equation, eq. (7). This defines the fraction of the deformation length that will be reserved for more rigid near-body deformation. If this is not set then the value of α will be computed dynamically based on equation (9). Usually the dynamic alpha provides better results than providing a constant value with this setting, so it is advised not to adjust this parameter.

5. `gridMotionAlphaFactor`

This variable is the η parameter in the dynamic alpha equation (9). The dynamic alpha will increase the stiff region around bodies when they have more complex deformations. The default value is set to 5 which appears to be sufficient for most applications.

6. `gridMotionAlphaFloor`

This variable sets the minimum value for α that will be used when the dynamic alpha computation is being used. This provides a guarantee that near wall regions will remain rigid and orthogonal. The default setting for this value is 0.1.

7. `gridMoveErrorTol`

This variable sets the error tolerance of the approximation algorithm. It is a percentage of the minimum edge length attached to the node. The default value is 0.1 which is usually sufficient for most applications. A larger value will degrade mesh quality but may improve speed, while a smaller value can improve mesh quality if approximation errors were the cause. Generally this value does not need to be changed from the default setting.

8. `gridDeformationSymmetry`

This variable is set if you have a symmetry plane at $x = 0$, $y = 0$, or $z = 0$. It is set to the string "x", "y", or "z" respectively to support the symmetry plane motion. When using a symmetry plane, the plane itself will be allowed to float and the deformations will be reflected so that the interpolation at the symmetry plane does not include out of plane deflections.

9. `gridDeformationPlane`

This variable is set for 2-D simulations to prevent deformations in the extruded direction. It is set to a string specifying the extruded axis.

In addition to the global parameters described earlier there are controls that can be provided on a per boundary basis. Note, if any boundary is not marked with a descriptor for mesh motion, then those nodes will be deformed just as the interior nodes are deformed. The boundary condition settings for the grid mover are as follows:

1. `fixed`

This setting tells the grid mover that this surface is not moving and so the deformation is identically zero for these nodes.

2. `moving`

This setting tells the grid mover that this surface will be receiving displacements from another module and as such the surface will be moving.

3. constrainedMotion

This setting tells the grid mover that the nodes on this surface will be allowed to move but that they are constrained to move within a specified surface. The specified surface can be a plane, infinite cylinder, or a discrete geometry file. Constrained motion surfaces are used to interface between moving surfaces and fixed surfaces. Examples of setting a plane constrained motion boundary would which constrains to the plane $x = 0$ defined using a point and normal would follow an input such as:

```
back=viscousWall(adiabatic,constrainedMotion=
                plane(point=[0,0,0],normal=[1,0,0])),
```

A cylinder is similarly defined using a cylinder key word given two points on the cylinder axis and the radius, which would be defined using an input similar to:

```
case=viscousWall(adiabatic,constrainedMotion=
                cylinder(p1=[0,0,0],p2=[4,0,0],radius=.5)),
```

Finally a general geometry constraint can be specified using a geometry file using a **geoFile** keyword such as:

```
case=viscousWall(adiabatic,
                constrainedMotion=geoFile(file="case.geom")),
```

The constrained motion option also allows intersecting constrained surfaces in which case the points that share both surfaces will be constrained to move along the curve of intersection.

4. deformWeight

By default, each node is assigned a weight proportional to its local area which means that a large surface may have undue influence over the deformation of grid nodes around a nearby smaller surface. The **deformWeight** can be added to a boundary to increase (if greater than 1) or decrease (if less than 1) the influence of that boundary surface on the volume deformation.

5. farfieldDeform

Since farfield boundaries usually have very large areas compared to the interior surfaces it is often helpful to mark the farfield boundaries with the **farfieldDeform** flag. This flag will automatically provide a **deformWeight** to the farfield boundaries that is formed by the ratio of the all non-farfield boundaries to the farfield area. This will reduce the influence of the farfield on the volume deformation.

APPENDIX

A Governing Equations

A finite-volume procedure is applied to discretize the flow equations. After integration over a computational cell, the governing equations for a three-dimensional flow with non-equilibrium chemistry and equilibrium internal energy, written in vector form for an arbitrary control volume Ω_c (closed by a boundary $\partial\Omega_c$) are:

$$\frac{d}{dt} \int_{\Omega_c(t)} Q \, dV + \int_{\partial\Omega_c(t)} (F_i - F_v) \, dS = 0 \quad (14)$$

where the vectors of conservative state variables, Q , inviscid flux, F_i , viscous flux, F_v , and chemistry source term, \dot{W} , are given by:

$$Q = \begin{bmatrix} \rho \\ \rho \tilde{u} \\ \rho e_0 \end{bmatrix}, \quad F_i = \begin{bmatrix} \rho \tilde{u} \cdot \tilde{n} \\ (\rho \tilde{u} \tilde{u} + p \tilde{I}) \cdot \tilde{n} \\ (\rho e_0 + p) \tilde{u} \cdot \tilde{n} \end{bmatrix}, \quad F_v = \begin{bmatrix} 0 \\ \tilde{\tau} \cdot \tilde{n} \\ (\tilde{u} \cdot \tilde{\tau} + \tilde{q}) \end{bmatrix}. \quad (15)$$

A.1 Thermodynamic Model

The equations are closed using the ideal gas model where pressure is defined by

$$p = \rho \tilde{R} T. \quad (16)$$

In Eq. (16), pressure is related to gas temperature, which is determined from the internal energy, $e_{internal}$. Moreover,

$$e_0 = \frac{1}{2} \tilde{u} \cdot \tilde{u} + e_{internal}, \quad (17)$$

where the total energy e_0 appears in the vector of conserved variables.

$$e_{internal} = \frac{\tilde{R} T}{\gamma - 1} \quad (18)$$

Note that $\tilde{R} = \hat{R}/m$ is the specific gas constant which is the ratio of the universal gas constant \hat{R} to the mean molecular mass, m .

B Formulation for Axisymmetric Flows with Swirl

The axisymmetric flow formulation follows the same finite volume discretization procedures that are used for three dimensional flows, except that the formulation begins with the governing equations written in cylindrical coordinates (*i.e.*, in (z, r, θ) coordinates which correspond to (x, y, z)). The governing equations in cylindrical coordinates are²:

$$\frac{\partial Q}{\partial t} + \frac{\partial E}{\partial z} + \frac{1}{r} \frac{\partial(rF)}{\partial r} + \frac{1}{r} \frac{\partial G}{\partial \theta} + H = \frac{\partial E_v}{\partial z} + \frac{1}{r} \frac{\partial(rF_v)}{\partial r} + \frac{1}{r} \frac{\partial G_v}{\partial \theta} + H_v \quad (19)$$

where the vector of conservative state variables, Q , inviscid fluxes, E , F , G , viscous fluxes, E_v , F_v , G_v and coordinate transformation source terms, H and H_v , are given by:

²For simplicity, only the nonreacting equations are discussed here.

$$Q = \begin{bmatrix} \rho \\ \rho u_z \\ \rho u_r \\ \rho u_\theta \\ \rho e_0 \end{bmatrix}, \quad (20)$$

$$E = \begin{bmatrix} \rho u_z \\ \rho u_z^2 + p \\ \rho u_z u_r \\ \rho u_z u_\theta \\ (\rho e_0 + p) u_z \end{bmatrix}, \quad F = \begin{bmatrix} \rho u_r \\ \rho u_r u_z \\ \rho u_r^2 + p \\ \rho u_r u_\theta \\ (\rho e_0 + p) u_r \end{bmatrix}, \quad G = \begin{bmatrix} \rho u_\theta \\ \rho u_z u_\theta \\ \rho u_r u_\theta \\ \rho u_\theta^2 + p \\ (\rho e_0 + p) u_\theta \end{bmatrix}, \quad (21)$$

$$E_v = \begin{bmatrix} 0 \\ \tau_{zz} \\ \tau_{zr} \\ \tau_{z\theta} \\ u_z \tau_{zz} + u_r \tau_{zr} + u_\theta \tau_{z\theta} - q_z \end{bmatrix}, \quad (22)$$

$$F_v = \begin{bmatrix} 0 \\ \tau_{rz} \\ \tau_{rr} \\ \tau_{r\theta} \\ u_z \tau_{rz} + u_r \tau_{rr} + u_\theta \tau_{r\theta} - q_r \end{bmatrix}, \quad (23)$$

$$G_v = \begin{bmatrix} 0 \\ \tau_{\theta z} \\ \tau_{\theta r} \\ \tau_{\theta\theta} \\ u_z \tau_{\theta z} + u_r \tau_{\theta r} + u_\theta \tau_{\theta\theta} - q_\theta \end{bmatrix}, \quad (24)$$

$$H = -\frac{1}{r} \begin{bmatrix} 0 \\ 0 \\ p + \rho u_\theta^2 \\ 0 \\ 0 \end{bmatrix}, \quad H_v = \frac{1}{r} \begin{bmatrix} 0 \\ 0 \\ -\tau_{\theta\theta} \\ \tau_{r\theta} \\ 0 \end{bmatrix}. \quad (25)$$

The shear stress components are defined as:

$$\begin{aligned}
\tau_{zz} &= 2\mu_T \frac{\partial u_z}{\partial z} + \lambda \nabla \cdot \vec{V} \\
\tau_{rr} &= 2\mu_T \frac{\partial u_r}{\partial r} + \lambda \nabla \cdot \vec{V} \\
\tau_{\theta\theta} &= 2\mu_T \left(\frac{1}{r} \frac{\partial u_\theta}{\partial \theta} + \frac{u_r}{r} \right) + \lambda \nabla \cdot \vec{V} \\
\tau_{r\theta} = \tau_{\theta r} &= \mu_T \left[r \frac{\partial}{\partial r} \left(\frac{u_\theta}{r} \right) + \frac{1}{r} \frac{\partial u_r}{\partial \theta} \right] \\
\tau_{\theta z} = \tau_{z\theta} &= \mu_T \left[\frac{\partial u_\theta}{\partial z} + \frac{1}{r} \frac{\partial u_z}{\partial \theta} \right] \\
\tau_{zr} = \tau_{rz} &= \mu_T \left[\frac{\partial u_z}{\partial r} + \frac{\partial u_r}{\partial z} \right],
\end{aligned} \tag{26}$$

where $\nabla \cdot \vec{V} = \frac{1}{r} \frac{\partial(ru_r)}{\partial r} + \frac{1}{r} \frac{\partial u_\theta}{\partial \theta} + \frac{\partial u_z}{\partial z}$ and $\lambda = -\frac{2}{3}\mu_T$. The heat flux vector is defined as:

$$q_z = -k_T \frac{\partial T}{\partial z}, \quad q_r = -k_T \frac{\partial T}{\partial r}, \quad q_\theta = -\frac{k_T}{r} \frac{\partial T}{\partial \theta}, \tag{27}$$

and the total viscosity and thermal conductivity coefficients are:

$$\mu_T = \mu + \mu_{turb}, \quad k_T = k + k_{turb}. \tag{28}$$

The axisymmetric equations with swirl are found by setting the θ derivatives to zero and by identifying the cylindrical velocity components (u_z, u_r, u_θ) with the Cartesian components (u, v, w) . Furthermore, the equations are multiplied through by $2\pi r$, and the areas and volumes are adjusted to account for a 360° rotation about the x -axis. A grid may thus be translated by an arbitrary amount in the z direction, and the axisymmetric flow will be computed in the $x-y$ plane.

The resulting equations are:

$$\frac{\partial(2\pi r Q)}{\partial t} + \frac{\partial(2\pi r E)}{\partial z} + \frac{\partial(2\pi r F)}{\partial r} + 2\pi r H = \frac{\partial(2\pi r E_v)}{\partial z} + \frac{\partial(2\pi r F_v)}{\partial r} + 2\pi r H_v, \tag{29}$$

and the 2D/axisymmetric finite volume form is:

$$\frac{d}{dt} \int Q dS' + \int [(E - E_v)\hat{n}_x + (F - F_v)\hat{n}_y] dl' = - \int (H - H_v) dS'_y = - \int r(H - H_v) dl'_y. \tag{30}$$

The quantity dS' is the modified cell area ($dS' = 2\pi r dS$), and dl' is the modified side length of each cell ($dl' = 2\pi r dl$). The factor $2\pi r$ is retained so that the integrations are taken over a 360° slice that is rotated about the x -axis³. The final step of identifying $dS'_y = r dl'_y$ is to ensure that a constant pressure uniform flow is numerically preserved.

³In the actual implementation, exact analytic area and volume integrations are used rather than the second order approximations given here.

C Turbulence Models

Several turbulence models are implemented in CHEM, including the algebraic Baldwin-Lomax[5], the one-equation Spalart-Allmaras[6], and a family of two-equation models including the SST formulation by Menter[7]. These models are briefly described in the following sections.

C.1 Spalart-Allmaras Model

The defining equations for this model are written as follows:

Kinematic Eddy Viscosity: $\nu_t = \tilde{\nu} f_{v1}$.

Eddy Viscosity Equation:

$$\frac{\partial \rho \tilde{\nu}}{\partial t} + \frac{\partial u_j \rho \tilde{\nu}}{\partial x_j} = \rho c_{b1} \tilde{S} \tilde{\nu} - \rho c_{w1} f_w \left(\frac{\tilde{\nu}}{y} \right)^2 + \frac{\rho}{\sigma} \frac{\partial}{\partial x_k} \left[(\nu + \tilde{\nu}) \frac{\partial \tilde{\nu}}{\partial x_k} \right] + \frac{\rho c_{b2}}{\sigma} \frac{\partial \tilde{\nu}}{\partial x_k} \frac{\partial \tilde{\nu}}{\partial x_k}, \quad (31)$$

where the last two terms on the left hand side represent turbulent diffusion, and tensor notation is employed (repeated indexes j, k denote summations). The first term on the right is the turbulence production, while the second term denotes the destruction due to the presence of a wall.

Closure Coefficients and Auxiliary Relations:

$$\begin{aligned} c_{b1} &= 0.1355, \quad c_{b2} = 0.622, \quad c_{v1} = 7.1, \quad \sigma = 2/3, \\ c_{w1} &= \frac{c_{b1}}{\kappa^2} + \frac{(1 + c_{b2})}{\sigma}, \quad c_{w2} = 0.3, \quad c_{w3} = 2, \quad \kappa = 0.41, \\ f_{v1} &= \frac{\chi^3}{\chi^3 + c_{v1}^3}, \quad f_{v2} = 1 - \frac{\chi}{1 + \chi f_{v1}}, \\ f_w &= g \left[\frac{1 + c_{w3}^6}{g^6 + c_{w3}^6} \right]^{1/6}, \quad \chi = \frac{\tilde{\nu}}{\nu}, \quad g = r + c_{w2}(r^6 - r), \\ r &= \frac{\tilde{\nu}}{\tilde{S} \kappa^2 y^2}, \quad \tilde{S} = S + \frac{\tilde{\nu}}{\kappa^2 y^2} f_{v2}, \quad S = \sqrt{2 \Omega_{ij} \Omega_{ij}}. \end{aligned}$$

The tensor $\Omega_{ij} = \frac{1}{2}(\partial u_i / \partial x_j - \partial u_j / \partial x_i)$ is one half of the vorticity tensor, and y is the distance to the closest wall surface. For simplicity, no tripping term is included.

The value $\tilde{\nu}$ at the wall boundary is set to zero, and the value of ν_t in the freestream is selected as $\nu_t = 10^{-3} \nu$.

The corresponding integral form of Eq. (31) can be included into the system of governing equations, Eq. (14). The resulting integral equations are solved, in a decoupled form, using the same algorithmic strategies as applied in the mean flow equations. We do perform a simple transformation to make the integration of the second diffusion term in the equation more straightforward to evaluate. This equivalent form that transforms all diffusion operators to face operations is given as:

$$\rho \frac{D \tilde{\nu}}{Dt} = \rho c_{b1} \tilde{S} \tilde{\nu} - \rho c_{w1} f_w \left(\frac{\tilde{\nu}}{y} \right)^2 + \frac{\rho}{\sigma} \left\{ \nabla \cdot [(\nu + (1 + c_{b2}) \tilde{\nu}) \nabla \tilde{\nu}] - c_{b2} \tilde{\nu} \nabla \cdot (\nabla \tilde{\nu}) \right\} \quad (32)$$

C.2 Compressible Form of Spalart-Allmaras

There is also an alternative compressible form of the Spalart-Allmaras model as developed by Catris and Aupoix. In this version of the model, the diffusion quantity is taken to be $\sqrt{\rho} \tilde{\nu}$ giving rise to the equation:

$$\rho \frac{D\tilde{\nu}}{Dt} = \rho c_{b1} \tilde{S} \tilde{\nu} - \rho c_{w1} f_w \left(\frac{\tilde{\nu}}{y} \right)^2 + \frac{1}{\sigma} \left\{ \nabla \cdot (\mu \nabla \tilde{\nu}) + \nabla \cdot (\sqrt{\rho} \tilde{\nu} \nabla \sqrt{\rho} \tilde{\nu}) + c_{b2} (\nabla \sqrt{\rho} \tilde{\nu})^2 \right\}. \quad (33)$$

In this version of the model, all of the coefficients remain the same as the original Spalart-Allmaras model.

C.3 Baseline Model (BSL)

It is well known that two-equation eddy-viscosity *low-Reynolds-number* turbulence models are among the most widely used models for engineering applications today, and the $k-\epsilon$ model with damping functions near the wall is the most popular. However, the $k-\epsilon$ model often suffers from numerical stability problems due to disparate turbulent time scales. Another well-known two-equation turbulence model is the $k-\omega$ model, developed by Wilcox[8]. It has the advantage that it does not require damping functions in the viscous sublayer and that the equations are less stiff near the wall, therefore it is superior to $k-\epsilon$ model with regard to numerical stability. However, when applied to the free shear layers, it is found that there is a strong dependency of the results on the freestream value of ω [9, 7]. Menter created a new model, called baseline (BSL) model, by blending the $k-\epsilon$ and $k-\omega$ models[10]. It utilizes the $k-\omega$ model in the wall region and gradually switches to the $k-\epsilon$ model away from the wall. To achieve this, the $k-\epsilon$ model is first diffusion term is added (another diffusion term associated with turbulent kinetic energy is neglected in the formulation under certain assumptions[11]). The original $k-\omega$ equations are then multiplied by a blending function F_{bsl} , the transformed $k-\epsilon$ equations are multiplied by $(1 - F_{bsl})$, and then both are added together. The blending function F_{bsl} is designed so that it is unity at the wall, and gradually approaches zero away from the wall. Note that the $k-\omega$ model can be easily obtained by setting $F_{bsl} = 1$ identically. In order to accurately predict adverse pressure gradient flows, especially in the wake region, Menter[10] modified the BSL model by including the transport of the principal turbulent shear stress[12] in the eddy-viscosity formulations. This leads to the shear-stress transport (SST) model. In the present study, both BSL and SST models are discussed.

The defining equations for the BSL model are written as:

Kinematic Eddy Viscosity:

$$\nu_t = k/\omega, \quad (34)$$

Turbulent Stress Tensor:

$$\tau'_{ij} = \mu_t \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right) - \frac{2}{3} (\mu_t \nabla \cdot \tilde{u} + \rho k) \delta_{ij}, \quad (35)$$

$$i = 1, 2, 3, \quad j = 1, 2, 3,$$

Turbulent Kinetic Energy Equation:

$$\rho \frac{Dk}{Dt} = \tau'_{ij} \frac{\partial u_i}{\partial x_j} - \beta^* \rho \omega k + \frac{\partial}{\partial x_j} [(\mu + \mu_t \sigma_k) \frac{\partial k}{\partial x_j}], \quad (36)$$

Turbulent Dissipation Equation:

$$\rho \frac{D\omega}{Dt} = \frac{\gamma}{\nu_t} \tau'_{ij} \frac{\partial u_i}{\partial x_j} - \beta \rho \omega^2 + \frac{\partial}{\partial x_j} [(\mu + \mu_t \sigma_\omega) \frac{\partial \omega}{\partial x_j}] + 2(1 - F_{bsl}) \rho \sigma_{\omega^2} \frac{1}{\omega} \frac{\partial k}{\partial x_j} \frac{\partial \omega}{\partial x_j}. \quad (37)$$

Closure Coefficients:

All the constants ϕ of the model are computed by blending the appropriate $k - \omega$ and $k - \epsilon$ constants, as follows:

$$\phi = F_{bsl}\phi_1 + (1 - F_{bsl})\phi_2, \quad (38)$$

where the constants ϕ_1 ($k - \omega$) are:

$$\begin{aligned} \sigma_{k1} &= 0.5, & \sigma_{\omega1} &= 0.5, & \beta_1 &= 0.075, \\ \beta^* &= 0.09, & \kappa &= 0.41, & \gamma_1 &= \beta_1/\beta^* - \sigma_{\omega1}\kappa^2/\sqrt{\beta^*}, \end{aligned} \quad (39)$$

and the constants ϕ_2 ($k - \epsilon$) are:

$$\begin{aligned} \sigma_{k2} &= 1.0, & \sigma_{\omega2} &= 0.856, & \beta_2 &= 0.0828, \\ \beta^* &= 0.09, & \kappa &= 0.41, & \gamma_2 &= \beta_2/\beta^* - \sigma_{\omega2}\kappa^2/\sqrt{\beta^*}. \end{aligned} \quad (40)$$

The blending function F_{bsl} is defined as follows:

$$F_{bsl} = \tanh(\arg_{bsl}^4), \quad (41)$$

where

$$\arg_{bsl} = \min\left[\max\left(\frac{\sqrt{k}}{0.09\omega y}, \frac{500\nu}{y^2\omega}\right), \frac{4\rho\sigma_{\omega2}k}{CD_{k\omega}y^2}\right], \quad (42)$$

and y is the distance to the closest point away from the wall surface. In the above, $CD_{k\omega}$ is defined as:

$$CD_{k\omega} = \max\left(2\rho\sigma_{\omega2}\frac{1}{\omega}\frac{\partial k}{\partial x_j}\frac{\partial \omega}{\partial x_j}, 10^{-20}\right). \quad (43)$$

The boundary conditions for k and ω at a solid wall are:

$$k = 0, \quad \omega = 10\frac{6\nu}{\beta_1(\Delta y_1)^2}, \quad (44)$$

where Δy_1 is the distance from the first cell center to the solid wall. For rough walls, the user-specified equivalent sand grain roughness height (k_s) is used to compute the solid wall value of ω from:

$$\omega = \frac{u_\tau^2}{\nu} S_R \quad (45)$$

where

$$S_R = \begin{cases} \left(\frac{200}{k_s^+}\right)^2, & k_s^+ \leq 5 \\ \frac{100}{k_s^+} + \left[\left(\frac{200}{k_s^+}\right)^2 - \frac{100}{k_s^+}\right] e^{5-k_s^+}, & k_s^+ > 5 \end{cases} \quad (46)$$

A surface is considered to be hydraulically smooth⁴ when $k_s^+ = u_\tau k_s/\nu < 5$. A slightly rough wall boundary condition for ω can be derived by combining Eqs. (45) and (46) and to obtain[13]:

$$\omega = \frac{40000\nu}{k_s^2}, \quad (47)$$

in which k_s must be chosen to ensure that $k_s^+ < 5$. Equations (44) and (45) are equivalent when $\Delta y_1 = 0.1414k_s$.

⁴A hydraulically smooth surface is one in which the roughness height is smaller than the laminar sublayer.

The corresponding integral form of Eqs. (36) and (37) can be included into the system of governing equations, Eq. (14), by adding the following additional terms to the vectors Q , F , F_v , and W :

$$\begin{aligned} Q &= \begin{bmatrix} \rho k \\ \rho \omega \end{bmatrix}, F = \begin{bmatrix} \rho k \tilde{u} \cdot \tilde{n} \\ \rho \omega \tilde{u} \cdot \tilde{n} \end{bmatrix}, F_v = \begin{bmatrix} (\mu + \mu_t \sigma_k) \nabla k \cdot \tilde{n} \\ (\mu + \mu_t \sigma_\omega) \nabla \omega \cdot \tilde{n} \end{bmatrix}, \\ W &= \rho \begin{bmatrix} \tau'_{ij} \frac{\partial u_i}{\partial x_j} - \beta^* \rho \omega k \\ \frac{\gamma}{\nu_t} \tau'_{ij} \frac{\partial u_i}{\partial x_j} - \beta \rho \omega^2 + 2(1 - F_{bsl}) \rho \sigma_\omega \frac{1}{\omega} \frac{\partial k}{\partial x_j} \frac{\partial \omega}{\partial x_j} \end{bmatrix}. \end{aligned} \quad (48)$$

C.4 Shear Stress Transport Model (SST)

The SST model is similar to the BSL model described above, except that $\sigma_{k1} = 0.85$, and the eddy viscosity is defined as:

$$\nu_t = \frac{a_1 k}{\max(a_1 \omega, \Omega F_{sst})}, \quad (49)$$

where Ω is the absolute value of the vorticity, and $a_1 = 0.31$. The blending function F_{sst} is given by:

$$F_{sst} = \tanh(\arg_{sst}^2), \quad (50)$$

where

$$\arg_{sst} = \max\left(2 \frac{\sqrt{k}}{0.09 \omega y}, \frac{500 \nu}{y^2 \omega}\right). \quad (51)$$

C.5 The Wilcox (1998) $k - \omega$ Turbulence Model

The Wilcox (1998) two equation turbulence model[11] is an updated version of the 1988 model[8]. The main differences between the new version and the previous 1988 version currently in the CHEM code are in the coefficients of the dissipation terms, which now depend on local flow conditions, and in the values of two modeling constants. The complete turbulence model specification is given below.

Eddy Viscosity:

$$\mu_T = \rho k / \omega \quad (52)$$

Turbulence Kinetic Energy:

$$\frac{\partial \rho k}{\partial t} + \frac{\partial \rho U_j k}{\partial x_j} = \tau_{ij} \frac{\partial U_i}{\partial x_j} - \rho \beta^* k \omega + \frac{\partial}{\partial x_j} \left[(\mu + \sigma^* \mu_T) \frac{\partial k}{\partial x_j} \right] \quad (53)$$

Specific Dissipation Rate:

$$\frac{\partial \rho \omega}{\partial t} + \frac{\partial \rho U_j \omega}{\partial x_j} = \alpha \frac{\omega}{k} \tau_{ij} \frac{\partial U_i}{\partial x_j} - \rho \beta \omega^2 + \frac{\partial}{\partial x_j} \left[(\mu + \sigma \mu_T) \frac{\partial \omega}{\partial x_j} \right] \quad (54)$$

Closure Coefficients and Auxiliary Relations:

$$\begin{aligned}
\alpha &= \frac{13}{25}, \quad \beta = \beta_0 f_\beta, \quad \beta^* = \beta_0^* f_{\beta^*}, \quad \sigma = 0.5, \quad \sigma^* = 0.5 \\
\beta_0 &= \frac{9}{125}, \quad f_\beta = \frac{1 + 70\chi_\omega}{1 + 80\chi_\omega}, \quad \chi_\omega \equiv \left| \frac{\Omega_{ij}\Omega_{jk}S_{ki}}{(\beta_0^*\omega)^3} \right| \\
\beta_0^* &= \frac{9}{100}, \quad f_{\beta^*} = \begin{cases} 1 & : \chi_k \leq 0 \\ \frac{1+680\chi_k^2}{1+400\chi_k^2} & : \chi_k > 0 \end{cases}, \quad \chi_k \equiv \frac{1}{\omega^3} \frac{\partial k}{\partial x_j} \frac{\partial \omega}{\partial x_j}
\end{aligned} \tag{55}$$

$$\epsilon = \beta^* \omega k \quad \text{and} \quad l = k^{1/2}/\omega \tag{56}$$

The mean rotation and strain rate tensors are defined as:

$$\Omega_{ij} = \frac{1}{2} \left(\frac{\partial U_i}{\partial x_j} - \frac{\partial U_j}{\partial x_i} \right), \quad S_{ij} = \frac{1}{2} \left(\frac{\partial U_i}{\partial x_j} + \frac{\partial U_j}{\partial x_i} \right) \tag{57}$$

and the Reynolds stress tensor is:

$$\tau_{ij} = 2\mu_T S_{ij} - \frac{2}{3} \rho k \delta_{ij} \tag{58}$$

The implementation of the Wilcox 1998 turbulence model into CHEM is straightforward and requires only the modification of existing source terms and the definition of new constants. In addition, the computation of the coefficients in Eqs. 55-56 was added. The new turbulence model is invoked from the CHEM input file as `turbulence_model: Wilcox98`. All of the other turbulence model options, including the choices of compressibility correction and the production function, work with this new model.

C.6 Hybrid RANS/LES Model (LES)

The hybrid RANS/LES turbulence model[14] is an implementation of a multiscale turbulence model in which the eddy viscosity is a function of two turbulent length scales as opposed to just one for the previously discussed models. The basic idea of the hybrid model is that the largest turbulent scales are resolved on the computational mesh while the smallest, unresolved scales continue to be modeled. This requires the definition of appropriate length scales, a filtering mechanism to determine which scales are modeled and which are resolved, and a blending function to smoothly match the eddy viscosity between the two regimes.

The turbulent length scale is defined as:

$$L_T = \max(6.0 \sqrt{\frac{\nu_{tRANS}}{\Omega}}, l_T), \tag{59}$$

where Ω is the local mean flow vorticity that helps define an algebraic length scale, and l_T is a turbulent length scale associated with two equation RANS models. The subscript RANS indicates that the values are from the unfiltered RANS model, while the factor 6.0 is used by Nichols and Nelson to make the two length scales approximately equal for a simple test problem. Their definition of l_T is

$$l_T = k_{RANS}^{3/2} / \epsilon_{RANS}, \tag{60}$$

whereas the usual definition is

$$l_T = \beta^* k_{RANS}^{3/2} / \epsilon_{RANS} = k^{1/2} / \omega. \tag{61}$$

The latter value has been used in the CHEM code, which suggests that further tuning of the constant 6.0 in Eq. 59 may be necessary.

When the grid is substantially refined outside of the boundary layer, additional turbulent scales will be resolved, and the eddy viscosity will be decreased compared to its usual RANS value. The subgrid value for the turbulence kinetic energy is thus taken as a fraction of the RANS value:

$$k_{LES} = k_{RANS} f_d, \quad (62)$$

where the damping function f_d is defined as:

$$f_d = \frac{1}{2} \{1 + \tanh[2\pi(\Lambda - 0.5)]\}. \quad (63)$$

The function Λ determines which scales are locally resolved by comparing the turbulent length scale L_T to the local grid size:

$$\Lambda = \frac{1}{1 + \left(\frac{L_T}{2L_G}\right)^{4/3}}, \quad (64)$$

where

$$L_G = \max(\Delta x, \Delta y, \Delta z) \quad (65)$$

for 3D flows. For axisymmetric or 2D flows, the user-specified out-of-plane direction is ignored (see below for the `.vars` file specification).

Finally, the blended eddy viscosity is computed from

$$\nu_t = \nu_{tRANS} f_d + (1 - f_d) \nu_{tLES}, \quad (66)$$

where the LES-based subgrid eddy viscosity is given by

$$\nu_{tLES} = \min(0.0854 L_G \sqrt{k_{LES}}, \nu_{tRANS}). \quad (67)$$

By examining Eqs. 63-64, we find the following delineations between the RANS and LES modes:

RANS Mode

$$L_T \ll L_G : \Lambda \rightarrow 1 : f_d \rightarrow 1 \quad (68)$$

LES Mode

$$L_T \gg L_G : \Lambda \rightarrow 0 : f_d \rightarrow 0 \quad (69)$$

In other words, when the turbulent scales (L_T) are much smaller than local grid scale (L_G), implying that they cannot be adequately resolved, the RANS mode is used in the usual single scale turbulence model approach. On the other hand, when the turbulent scales are much larger than the local grid scale and can be resolved, the model switches smoothly to the LES mode. This results in a smaller eddy viscosity which is necessary to resolve the unsteady turbulent fluid motion that was originally damped out by the larger RANS eddy viscosity.

The above hybrid RANS/LES model must be associated with a two equation turbulence model, and thus the only addition to the `.vars` file for a hybrid RANS/LES turbulent run is the following line:

multi_scale: LES

Note, however, that since CHEM is a 3D code, the default option for the local grid scale (Eq. 65) is 3D. In order to make axisymmetric or 2D runs, the user must specify which out-of-plane direction is to be ignored when computing the local grid scale. This is done by replacing LES with either LES2DX, LES2DY, or LES2DZ, depending on which coordinate direction is to be ignored. For example, for a run in the xy plane, the z direction is out-of-plane and thus should be ignored, so we would use:

multi_scale: LES2DZ

to obtain the proper length scale calculation.

C.7 Turbulence Compressibility Corrections

Compressibility corrections for high-speed shear flows [15, 16] are implemented in the turbulence equations. These corrections are expected to improve the results on problems involving high speed free shear flows. For the BSL and SST turbulence models we implement Wilcox's [11] modification to the correction. The compressibility correction is implemented as a modification to the β and β^* destruction terms that appear in equations (36) and (37) such that the new constants are given by:

$$\beta_{corr}^* = \beta^*[1 + \xi^* F(M_t)], \quad (70)$$

$$\beta_{corr} = \beta - \beta^* \xi^* F(M_t), \quad (71)$$

where $F(M_t) = \max(M_t^2 - M_{t_0}^2, 0)$. We recover Sarkar's original model by setting $\xi^* = 1$ and $M_{t_0} = 0$, and Wilcox's modification by setting $\xi^* = \frac{3}{2}$, and $M_{t_0} = \frac{1}{4}$. The correction is disabled by setting $\xi^* = 0$.

C.8 The Dynamic Hybrid RANS-LES (DHRL) turbulence model

One hybridized RANS-LES modeling option provided in the *flowPsi* code is based on an averaging consistent framework that can couple any arbitrary RANS turbulence model with any arbitrary RANS model. In this model the RANS turbulence model is solved on the time averaged flowfield. In this framework there is always a consistent RANS interpretation in all regions, even those that are using the LES model. To derive the DHRL approach, the instantaneous velocity (u_i) is decomposed into Favre-averaged \bar{u}_i , resolved fluctuating (u''), and unresolved fluctuating (u'_i) components, *i.e.*:

$$\bar{u}_i = \langle \hat{\rho} \tilde{u}_i \rangle / \langle \hat{\rho} \rangle \quad (72)$$

$$u_i'' = \tilde{u}_i - \bar{u}_i \quad (73)$$

$$u_i' = u_i - \tilde{u}_i \quad (74)$$

Here the angle brackets denote Reynolds averaging. For stationary flow, this is equivalent to an infinite-time average.

The assumptions used in the derivation of the DHRL method are: (1) negligible correlation between the resolved and unresolved fluctuations, (2) scale similarity, *i.e.* $\hat{\rho} \tilde{u}_i \tilde{u}_j - \hat{\rho} \tilde{u}_i \tilde{u}_j = \alpha \tau_{ij}^{SGS}$ and $\hat{\rho} \tilde{u}_i' u_j' = \beta \hat{\rho} \bar{u}_i' \bar{u}_j'$, and (3) that model constants α and β are complementary. This yields the following blending method for RANS and LES stresses:

$$\tau_{ij} = \alpha \tau_{ij}^{SGS} + (1 - \alpha) \tau_{ij}^{RANS}, \quad (75)$$

$$\alpha = \underbrace{-\langle \hat{\rho} \rangle \overline{u_i'' u_j'' S_{ij}}}_{\text{Resolved Production}} / \left(\underbrace{\tau_{ij}^{RANS} \bar{S}_{ij}}_{\text{RANS Production}} - \underbrace{\bar{\tau}_{ij}^{SGS} \bar{S}_{ij}}_{\text{Inhomogeneous Production}} \right) \quad (76)$$

where τ_{ij}^{SGS} is the subgrid stress predicted by any candidate LES model (for MILES this is identically zero), and τ_{ij}^{RANS} is the turbulent stress predicted by any candidate RANS model. Refer to Bhushan and Walters[17] and Walters *et al.* [18] for further details.

The above approach is similar to dynamic LES model coefficient evaluation, where the secondary filter is in fact the Reynolds (or Favre) averaging operation. The numerator in Eq. (76)

represents the sum of production of turbulent kinetic energy (k) due to the resolved turbulent scales in the flow, and $\bar{\tau}_{ij}^{SGS}\bar{S}_{ij}$ which is the mean component of the subgrid scale turbulent kinetic energy production. The term in the denominator, $\tau_{ij}^{RANS}\bar{S}_{ij}$, is the production of k predicted by the RANS model. A similar expression to Eq. (75) is adopted for the turbulent heat flux:

$$q_j = \alpha q_j^{SGS} + (1 - \alpha) q_j^{RANS} \quad (77)$$

Equation (76) indicates that the model operates in a pure LES mode only if the resolved scale production is equal to the predicted RANS production; otherwise, the model behaves in a transitional mode where an additional RANS stress compensates for reduced LES content. This leads to a smooth variation of turbulent production across the transition region. In regions with zero LES content, i.e. numerically steady flow, the model operates in a pure RANS mode. Note that, according to Eq. (76), the value of α can become negative and/or singular. In practice, the value of α is limited to lie between 0 (pure RANS) and 1 (pure LES).

One further critical aspect of the DHRL method is that the RANS model terms are computed based solely on the Favre-averaged flowfield. For example, for a linear eddy viscosity model the turbulent stress and heat flux are computed as:

$$\tau_{ij}^{RANS} = \frac{2}{3} \rho k \delta_{ij} - \mu_T \left(\frac{\partial \bar{u}_i}{\partial x_j} + \frac{\partial \bar{u}_j}{\partial x_i} - \frac{2}{3} \frac{\partial \bar{u}_k}{\partial x_k} \delta_{ij} \right) \quad (78)$$

$$q_j^{RANS} = -k_T \frac{\partial \bar{T}}{\partial x_j} \quad (79)$$

Similarly, all of the terms in the model transport equations for turbulence model dependent variables (e.g. k and ω) are computed in terms of \bar{u}_i , including convective and production terms. In stationary flows, for example, the velocity field used to compute all RANS terms is obtained from a running time-average. Other appropriate averaging methods can be adopted for non-stationary flows. For example, as discussed below, the DHRL model can be run using the exponential averaging option for online statistics, in which case the averaged terms (e.g. \bar{u}) strictly represent time-filtered, rather than infinite averaged, quantities.

C.8.1 Current DHRL Implementation

For the `flowpsi` implementation the DHRL model has been implemented and verification tests have been performed for use with the SST $k - \omega$ model for the RANS component and MILES for the LES component. Wall functions are not currently supported, solid boundaries should be modeled as `viscousWall` conditions and mesh resolution should be sufficient to resolve the mean flow features for accurate boundary layer prediction. To implement DHRL, load the following modules:

```
loadModule: KOmegaModel
loadModule: dhrl
```

C.8.2 DHRL Inputs

The following inputs are required in the `.vars` file to properly enable DHRL:

```
turbulence_model: SST
multi_scale: DHRL
```

C.8.3 Optional inputs for DHRL include:

`dhrl_source_terms: on`

comment: This input parameter can be used to disengage the inclusion of the RANS source terms in the governing equations by setting to “off”. This allows the simulation to be effectively run in a pure LES mode, while simultaneously obtaining values for the averaged variable field statistics, as well as the RANS model variable fields (k and ω) based on the time-averaged velocity field from the LES simulation.

`exponentialMean: off`

comment: The standard use of DHRL is for stationary turbulent flow. To use time-filtering instead of infinite averaging, input “on” for `exponentialMean`.

`meanFreq: 1000`

comment: For standard DHRL (`exponentialMean: off`), the value of `meanFreq` should be set to a very large number, i.e. greater than the maximum expected timesteps to be run in the simulation, in order to ensure that the averaging counter does not reset during the simulation. If time filtering is used, the value of `meanFreq` determines the filter width in terms of the number of time steps

`meanCountReset: NO DEFAULT VALUE`

comment: If this parameter is set to a non-zero integer value, the counter used for the time averaging will be reset to that value. Note that the values of the averaged variables themselves will not be changed.

C.8.4 DHRL Outputs

The following additional variables are available as outputs when running the DHRL model:

`alpha_dhrl`

comment: This is the scalar value of the RANS-to-LES blending parameter

C.8.5 Recommended DHRL Start Up Procedure

A converged DHRL solution requires infinite-time averaged quantities within the model formulation itself. For this reason, it is desirable to remove start up transients as quickly as possible from the flow statistics. The following is a recommended procedure for reducing the overall run time required for a DHRL simulation.

1. (Optional) Obtain a converged or nearly converged RANS solution. This may be useful if, for example, the initialization may lead to unstable behavior for a pure LES simulation.
2. Run the DHRL model using initial solution either prescribed or obtained from precursor RANS simulation, using the `dhrl_source_terms: off` input option. This is effectively a MILES simulation but with collection of averaged statistics as well as development of the steady k and ω solution field based on the time-averaged velocity field. The simulation need not be run to statistical convergence, but should be run sufficiently long that LES-like behavior is clearly observed and any large-scale features of the initial field are no longer apparent in the results.

3. Run the DHRL model using the `dhrl_source_terms: on` (default) option and setting `meanCountReset: #`. A recommended value for `meanCountReset` is 100. If the simulation proves to be unstable, this value can be increased.
4. Continue running DHRL, *making sure that for any restarts the meanCountReset variable is not included in the .vars file*. Monitor averaged quantities of interest (e.g. drag) to determine variation of key variables as a function of simulation time. Determination of statistical convergence is identical to that for pure LES.

C.8.6 Optimal settings for MILES

The preferred method for running the DHRL model is using a MILES (Monotone Implicit LES) approach for the subgrid model whereby the implicit dissipation provided by the upwind scheme is used as a subgrid model. Since the upwind operators provide a natural low pass filter, they provide dissipation characteristics that are consistent with the requirements of a good LES filter. However, the standard Roe upwinding scheme can add too much dissipation and so some sort of dissipation lowering technique will be needed for the best results. Currently the suggested method is to utilize the skew-symmetric flux options provided by *flowPsi*. The recommended settings for this are:

```
inviscidFlux: ssf
LDS_useUpwind: 0.20
```

This will use the 4th order skew symmetric flux blended with 20% of the standard upwind flux. This amount of upwinding has been found to have good performance as a MILES LES subgrid model. This skew symmetric based flux will not work reliably on tetrahedral or prismatic element meshes, however often reasonable results can be obtained by refining the mesh with `refmesh` which splits the tetrahedral into hexahedral elements.

D Numerical Formulation

The numerical model employed in this study is a finite-volume technique that supports generalized grids[19] (*generalized grids* are discretizations composed of arbitrary polyhedra, including tetrahedra, prisms, pyramids, and hexahedra). Since this formulation is based on cell centered integrations, any grid type can be expressed, including “hanging” nodes found in adaptive mesh refinement (AMR).

D.1 Numerical Approximations to Spatial Integrals

The numerical solution of the governing equations, Eq. (14), is obtained by applying the finite volume method. This approach is frequently used because it can guarantee that numerical truncation errors do not violate conservation properties. The numerical integration of Eq. (14) begins with approximations to volume and surface integrals. For the volume integrals a second-order midpoint rule is used. For example, the numerical integration of Q results in

$$\int_{\Omega_c(t)} Q(\tilde{x}, t) dV = Q_c(t) \mathcal{V}_c(t), \quad (80)$$

where $Q_c(t)$ is the value of Q at the centroid of cell c , and $\mathcal{V}_c(t)$ is defined by

$$\mathcal{V}_c(t) = \int_{\Omega_c(t)} dV. \quad (81)$$

The numerical integration of the surface integral in Eq. (14) is accomplished by summing the contributions of each of the \mathbf{NF} faces of cell c . Each individual contribution is again approximated using the midpoint rule. The flux function itself will require additional numerical treatment, and will be discussed in later sections. For now, assume that the flux can be approximated by a function, \hat{F} , of conservative values to the left and right of the face. Given this, the numerical integration of $F = F_i - F_v$ results in the following

$$\int_{\partial\Omega_c(t)} F dS = \sum_{f=1}^{\mathbf{NF}_c} \int_{\partial\Omega_{c,f}(t)} F dS \approx \sum_{f=1}^{\mathbf{NF}_c} \mathcal{A}_{c,f}(t) \hat{F}_f, \quad (82)$$

where the area of the face, $\mathcal{A}_{c,f}(t)$, is defined as

$$\mathcal{A}_{c,f}(t) = \int_{\partial\Omega_{c,f}(t)} dS. \quad (83)$$

At this point, Eq. (14) is numerically approximated by the equation

$$\frac{d}{dt} [\mathcal{V}_c(t) Q_c(t)] + \sum_{f=1}^{\mathbf{NF}_c} \mathcal{A}_{c,f}(t) \hat{F}_f = 0. \quad (84)$$

Notice that the differential term that remains in this equation applies to the product of volume and conservative state vector for the cell. However, the variable that is the objective of these calculations is $Q_c(t)$, not $\mathcal{V}_c(t) Q_c(t)$. This problem is solved by applying the chain rule:

$$\frac{d}{dt} Q_c(t) \mathcal{V}_c(t) = Q_c(t) \frac{d}{dt} \mathcal{V}_c(t) + \mathcal{V}_c(t) \frac{d}{dt} Q_c(t). \quad (85)$$

The derivative of volume with respect to time can be converted into a spatial integral through the use of an identity for integration over time-dependent domains[20]:

$$\begin{aligned} Q_c \frac{d}{dt} \mathcal{V}_c(t) &= Q_c \frac{d}{dt} \int_{\Omega_c(t)} dV \\ &= Q_c \int_{\partial\Omega_c(t)} \tilde{u}_\Omega \cdot \tilde{n} dS \\ &\approx Q_c \sum_{f=1}^{\mathbf{NF}_c} \mathcal{A}_{c,f}(t) (\tilde{u}_{\Omega,f} \cdot \tilde{n}_{c,f}). \end{aligned} \quad (86)$$

This equation, known as the geometric conservation law[21], is necessary for correct time integration when mesh deformation is present. Given this, the solution method can now be described in terms of a system of ordinary differential equations of the form

$$\mathcal{V}_c \frac{d}{dt} Q_c = R_c, \quad (87)$$

where R_c is given by the expression

$$R_c = \sum_{f=1}^{\mathbf{NF}_c} \mathcal{A}_{c,f} \hat{F}_{v,f} - \sum_{f=1}^{\mathbf{NF}_c} \mathcal{A}_{c,f} \hat{F}_f - Q_c \sum_{f=1}^{\mathbf{NF}_c} \mathcal{A}_{c,f} (\tilde{u}_{\Omega,f} \cdot \tilde{n}_{c,f}). \quad (88)$$

Eqs. (87) and (88) describe a system of ordinary differential equations that numerically model the time evolution of the fluid dynamics equations when simultaneously satisfied for all cells in

the mesh. To represent this fact, the cell subscript c will be dropped. Thus Q_c represents the fluid state for cell c , while Q represents the fluid states of all cells in the mesh. For example, while Eq. (87) represents the cell by cell differential equations, the global system of equations is given by

$$\mathcal{V} \frac{d}{dt} Q(t) = R(Q(t), t). \quad (89)$$

D.2 Geometric Integrations

In the previous section, numerical integrations over surfaces and volumes were described without explicitly defining the formulas used to evaluate Eqs. (81) and (83). Here the numerical integrations that approximate the volumes and areas of generalized cells and faces are discussed in some detail. Several considerations have to be made regarding these computations. First, when the nodes of a generalized face are not co-planar, the geometry of the face is not uniquely specified. A unique specification of the geometry of such faces can be obtained by subdividing the face into triangles: using a symmetric decomposition, each edge of the face, combined with the face centroid, creates a triangle. This strategy allows the face geometry to be independent of data-structures used to describe generalized meshes (the geometry is the same regardless of the ordering of the face edges). Using this technique, the area for the face is determined numerically using the following summation

$$\tilde{\mathcal{A}}_{c,f} = \frac{1}{2} \sum_{e=1}^{\mathbf{NE}_f} (\tilde{x}_{1,e} - \tilde{x}_{c,f}) \times (\tilde{x}_{2,e} - \tilde{x}_{c,f}), \quad (90)$$

where $\tilde{x}_{1,e}$ and $\tilde{x}_{2,e}$ are the positions of the two nodes of edge e in counter-clockwise order, and $\tilde{x}_{c,f}$ is the face centroid (approximated by an edge-length-weighted sum of edge centers). Here the computed area is a vector, which is represented as the product of the face normal vector and the face area, $\tilde{\mathcal{A}} = \mathcal{A}\tilde{n}$.

For numerical approximations of the volume integrals, it is essential to achieve a geometric conservation of volume, avoiding truncation errors in integrating volumes that would yield inaccurate total volume computations. For example, one would expect that the sum of the individual cell volumes would be equal to the total grid volume if exact arithmetic were used. To accomplish this, the volume of a cell is determined by decomposing it into tetrahedra and summing the individual tetrahedral volumes. The volume of a tetrahedron is given by

$$\mathcal{V}_{tet} = \frac{1}{3!} \left(\tilde{a} \cdot (\tilde{b} \times \tilde{c}) \right), \quad (91)$$

where \tilde{a} , \tilde{b} , and \tilde{c} are the three edge vectors from a given vertex of the tetrahedron. The volume of a generalized cell is then computed by using the triangulation of each face and the cell centroid: each surface triangle and the cell centroid define a tetrahedron. The cell volume is the sum of the tetrahedra volumes. Thus, the volume of a cell is computed by combining Eqs. (90) and (91), and factoring, to obtain

$$\mathcal{V}_c = -\frac{1}{3} \sum_{f=1}^{\mathbf{NF}_c} (\tilde{x}_{c,c} - \tilde{x}_{c,f}) \cdot \tilde{\mathcal{A}}_{c,f}, \quad (92)$$

where $\tilde{x}_{c,c}$ is the cell centroid (approximated by the area-weighted sum of face centroids). Notice that this computation assumes that the face normals point outward from the cell, which explains the negative sign. Using this approach, one can compute the correct volume of any cell, including highly non-convex cells.

D.3 Inviscid Flux Schemes

The default inviscid flux used in the *flowPsi* is Torro's HLLC scheme [22]. However, a low dissipation flux is also available when may be enabled by setting the `inviscidFlux` to either 'kec' or 'ssf'. The 'ssf' scheme will revert to a 4th order accurate skew symmetric scheme on Cartesian meshes.

D.4 Preconditioned HLLC flux option

The robust Chorin-Turkel preconditioning is implemented to facilitate efficient and accurate solution for low speed flows. This is particularly useful for RANS computations. Preconditioning provides two benefits for approximate Riemann solver codes. First, preconditioning rescales the acoustic wave speeds to be comparable in magnitude to the convective waves. This allows for improved iterative convergence to the steady state solution. Second, preconditioning reduces the dissipation of standard flux schemes such as HLLC. For flows below about Mach 0.2, the dissipation of the scheme can make accurate solution difficult or impossible.

D.4.1 Background on preconditioning

The basic idea behind preconditioning is the provision of modified time derivative which rescales the the acoustic wave speeds without changing the convective wave speeds. The time derivative for preconditioning is modified in the following way:

$$MP^{-1} \frac{\partial q}{\partial t} = \text{div}(F(q)) + \text{div}(F_v(q)) \quad (93)$$

where $q = [T, \mathbf{u}, p_g]^T$ is the primitive vector, M is the conservative transformation matrix and P^{-1} is the inverse of the preconditioning matrix. The conservative transformation matrix is given by

$$\frac{\partial Q}{\partial q} = M = \begin{bmatrix} -\frac{\rho}{T} & 0 & 0 & 0 & \frac{\rho}{RT} \\ -\frac{\rho u}{T} & \rho & 0 & 0 & \frac{\rho u}{RT} \\ -\frac{\rho v}{T} & 0 & \rho & 0 & \frac{\rho v}{RT} \\ -\frac{\rho w}{T} & 0 & 0 & \rho & \frac{\rho w}{RT} \\ f_T & \rho u & \rho v & \rho w & f_p \end{bmatrix}, \quad (94)$$

where f_T and f_p are the derivatives of ρe_0 with respect to pressure and temperature which is given as

$$f_T = -\frac{\rho h_0}{T} + \frac{\gamma}{\gamma - 1} \rho \tilde{R}, \quad (95)$$

and

$$f_p = \frac{h_0}{\tilde{R}T} - 1 \quad (96)$$

The preconditioning matrix is given by

$$\mathbf{P} = \begin{bmatrix} 1 & 0 & 0 & 0 & -\frac{(1-\eta_p)(\gamma-1)T}{a^2\rho} \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & \eta_p \end{bmatrix} \quad (97)$$

The inverse of \mathbf{P} can be found as

$$\mathbf{P}^{-1} = \begin{bmatrix} 1 & 0 & 0 & 0 & -\frac{(1-1/\eta_p)(\gamma-1)T}{a^2\rho} \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1/\eta_p \end{bmatrix}, \quad (98)$$

where η_p is the preconditioning parameter which is given by

$$\eta_p = M_r^2 / (1 - M_r^2) \quad (99)$$

where M_r is a local Mach number which is bounded such that $M_r^2 \leq \frac{1}{2}$. In addition, to prevent instability in the case caused by Mach number approaching zero, a lower bound on M_r is provided based on an estimated free stream mach number provided by the user.

The eigenvalues for this modified system of equations are given by

$$\Lambda = [u_n, u_n, u_n, \frac{u_n(1 + \eta_p) - \sigma}{2}, \frac{u_n(1 + \eta_p) + \sigma}{2}]. \quad (100)$$

where $\sigma = \sqrt{u_n^2(1 - \eta_p)^2 + 4\eta_p a^2}$.

The change in wave speeds by the preconditioning matrix changes the eigenvectors and eigenvalues associated with the approximate Riemann solver. In the case of the HLLC flux, it is sufficient to modify the eigenvalues used in the construction of the HLLC flux. For a Roe based scheme the new eigenvectors are used to construct a preconditioned Roe matrix. We use the modified HLLC scheme on the RHS of the equations, but use the preconditioned Roe matrix as an approximate linearization.

D.4.2 Skew Symmetric low dissipation flux schemes

Our goal is to develop a low dissipation numerical scheme for unstructured grids by forming a hybrid of a low dissipation skew symmetric flux designed for Cartesian meshes with traditional upwinded unstructured methods such that we recover the Cartesian mesh algorithm for regular meshes and also provide a low dissipation characteristic on irregular meshes. The Cartesian formulation in which this is based on are 2^{nd} and 4^{th} order fluxes from the skew-symmetric family of flux functions. For this discussion, consider the 2^{nd} order skew symmetric flux that is evaluated at cell faces as a function of averaged variables by the relation

$$F_{avg}(l, r) = \rho(p_{avg}, T_{avg}) (\vec{u}_{avg} \cdot \vec{n}) \begin{bmatrix} 1 \\ \vec{u}_{avg} \\ e(p_{avg}, T_{avg}) + k_{avg} \end{bmatrix} + \begin{bmatrix} 0 \\ p_{avg} \vec{n} \\ p_{avg} (\vec{u}_{avg} \cdot \vec{n}) \end{bmatrix} \quad (101)$$

where the density, $\rho(p, T)$, and internal energy, $e(p, T)$, are provided by the equation-of-state, \vec{n} is the face normal vector, and the averaged variables p_{avg} , T_{avg} , \vec{u}_{avg} , and k_{avg} are provided by the following averaging relations:

$$\begin{aligned} p_{avg} &= \frac{1}{2} (p_l + p_r) \\ T_{avg} &= \frac{1}{2} (T_l + T_r) \\ \vec{u}_{avg} &= \frac{1}{2} (\vec{u}_l + \vec{u}_r) \\ k_{avg} &= \vec{u}_{avg} \cdot \vec{u}_{avg} - \frac{1}{4} (\vec{u}_l \cdot \vec{u}_l + \vec{u}_r \cdot \vec{u}_r) \end{aligned} \quad (102)$$

For an ideal gas, the equations-of-state provide the following functions:

$$\begin{aligned}\rho(p, T) &= \frac{p}{\tilde{R}T}, \\ e(p, T) &= \frac{\tilde{R}}{\gamma - 1}T.\end{aligned}\tag{103}$$

The above formulation reproduces the compressible kinetic energy consistent (KEC) form of Subbareddy *et al.*[23] due to the definition of average fluid kinetic energy, k_{avg} . However, this scheme differs from the Subbareddy flux form in that averages of temperature and pressure are formed and then density and internal energy are derived from the equation-of-state (EoS), whereas the original KEC flux formed independent averages for these variables. Other options for averaging, including density and internal energy averaging, were also evaluated and it was found that the above averaging process generally produced errors that are lower by a factor of up to two.

The second order skew-symmetric flux of eq. (101) can be used as a basis to construct a 4th order scheme following the construction of Pirozzoli[24, 25] and given here as the form:

$$F_{SSF} = (1 + 2\beta)F_{avg}(l, r) - \beta(F_{avg}(ll, r) + F_{avg}(r, rr)),\tag{104}$$

where $\beta = \frac{1}{6}$ and ll denotes the value of the cell left-of-left and rr denotes the value of the cell right-of-right. This 4th order skew-symmetric flux (obtained when $\beta = \frac{1}{6}$) will be denoted as the SSF scheme. When $\beta = 0$ this flux reduces to the KEC scheme of Subbareddy *et al.*. This flux will provide the central flux scheme in our hybridized low dissipation scheme.

D.4.3 Unstructured Mesh Hybridization

For unstructured solvers we would like to modify the SSF scheme in such a way that it becomes identical to the 4th order scheme on Cartesian meshes but gracefully degrades to a low dissipation second order scheme on unstructured meshes. The flux written in equation (104) is not directly applicable to unstructured meshes because of the direct reference to left-of-left and right-of-right cells that have no clear definition in an irregular mesh structure. The general problem is illustrated in Figure 4 where \vec{x}_l and \vec{x}_r denote the cell centroids of the given polyhedral mesh, \vec{x}_f is the face centroid of a given polygonal face, while \vec{x}_c is the midpoint location between the cell faces where the centered flux is defined. Here \vec{x}_{ll} and \vec{x}_{rr} are the nominal locations of a Cartesian mesh overlayed onto this unstructured mesh. In general no cells are centered at these locations, but for a Cartesian mesh these points will correspond to the expected cell centers. The construction of an unstructured scheme requires the resolution of two issues: how to correct for the fact that the flux interpolation center (X_c) does not correspond to the face centroid (X_f) and how to obtain values for the left-of-left and right-of-right values in an unstructured context.

We can utilize the cell gradients computed on the left and right side to devise a correction for the offset between \vec{x}_f and \vec{x}_c . Note that the without this correction the scheme given in Eq. (104) will degrade to first order on unstructured grids unless mesh refinement is sufficiently smooth (e.g. $|\vec{x}_f - \vec{x}_c| = O(h^2)$). A correction that will restore second order accuracy for non-smooth refinements is accomplished by using the cell centered gradients at the left and right cells to shift the stencil such that it is centered on the face center. For a given interpolated value ϕ , the states interpolated to this new stencil are given by the correction

$$\begin{aligned}\phi'_l &= \phi_l + \nabla\phi_l \cdot (\vec{x}_f - \vec{x}_c) \\ \phi'_r &= \phi_r + \nabla\phi_r \cdot (\vec{x}_f - \vec{x}_c),\end{aligned}\tag{105}$$

which provides a second order correction provided that $\nabla\phi$ is computed to at least first order accuracy.




Figure 4: Geometry of flux stencil formation

To recover a 4th order accurate scheme on Cartesian meshes we need to establish a reconstruction for the left-of-left and right-of-right stencil locations that will correspond to the actual values at these cell locations on a Cartesian mesh, but are only functions of the available left and right values and gradients. To do this we start with the assumption that gradients will be computed at the cells using a least squares reconstruction. For this reconstruction on a Cartesian mesh, it can be show that the following relationships will identically hold:

$$\begin{aligned}\nabla\phi_l \cdot (\vec{x}_r - \vec{x}_l) &= \frac{\phi_r - \phi_l}{2} \\ \nabla\phi_r \cdot (\vec{x}_r - \vec{x}_l) &= \frac{\phi_{rr} - \phi_l}{2}\end{aligned}\tag{106}$$

Using these relations, we can then compute an effective left-of-left and right-of-right state using the cell values and their gradients. These are defined as follows:

$$\begin{aligned}\phi_{ll}' &= \phi_r' - 2\nabla\phi_l \cdot (\vec{x}_r - \vec{x}_l) \\ \phi_{rr}' &= \phi_l' + 2\nabla\phi_r \cdot (\vec{x}_r - \vec{x}_l)\end{aligned}\tag{107}$$

The approach produces exactly the 4th order flux on Cartesian grids and also provides a reasonable correction for general meshes. This approach has the advantage of only requiring the cell centered least squares gradient which will be required for the upwind flux as was observed in a similar formulation presented by Darwish and Moukalled[26]. Note, this implementation was arrived at after evaluating many subtle variations on the themes presented here. For example, a centered face interpolation using cubic polynomials is evaluated to replace the face averaged values for the skew-symmetric flux but did not produce 4th order convergence on Cartesian meshes due to the presence of aliased second order modes. Similarly, centered linear reconstructions at the face were evaluated to correct for location of the face center, but the resulting scheme is less stable and required more blended upwinding than the one presented here.

D.4.4 Hybrid Flux Scheme

In practice the skew symmetric fluxes (equation (104)) fail for shocked flows and therefore need to be hybridized with standard upwind fluxes to enable simulation of high speed flows. In addition, unstructured stencils can depart significantly from skew-symmetry and thus will require some stabilizing dissipation which is provided by blending to the upwind flux. This blending is achieved by using a simple weighted average of central flux and upwind flux schemes such that our final inviscid flux can be written as

$$F_i = (1 - \alpha_{diss})F_{SSF} + \alpha_{diss}F_{upwind}.\tag{108}$$

For this study we utilize the robust HLLC scheme[22] for the upwind flux. Since we are also using the upwind scheme to stabilize corrections on unstructured meshes we modify the upwind scheme to provide low dissipation in low Mach number regions by employing the simple scaling scheme proposed by Thornber *et al.*[27] whereby the left and right velocity states are modified according to the local face Mach number, M_f , according to the relations:

$$\begin{aligned}\vec{u}_{l,th} &= \frac{\vec{u}_l + \vec{u}_r}{2} + \min(M_f, 1) \frac{\vec{u}_l - \vec{u}_r}{2} \\ \vec{u}_{r,th} &= \frac{\vec{u}_l + \vec{u}_r}{2} + \min(M_f, 1) \frac{\vec{u}_r - \vec{u}_l}{2}\end{aligned}\tag{109}$$

Blending upwind fluxes will be motivated by three concerns: 1) geometric concerns where upwinding is used to stabilize the central flux schemes when skew-symmetry is lost, and 2) to provide a proper amount of added dissipation in regions with dissipative structures such as shocks and expansions, and 3) to provide a Monotone Implicit LES model (MILES) approach to subgrid filtering by employing a small amount of upwinding dissipation as an implicit subgrid model. Therefore α_{diss} will be computed based on mesh geometry (α_{geom}) and compressibility based (α_{comp}), and subgrid modeling based (α_{miles}) considerations. Ultimately α_{diss} will be determined by the maximum of these three factors. Since blending to upwinding generally introduces dissipation to the scheme, we would like to keep the non-physics based α_{geom} as small as possible. We use a heuristic approach to add dissipation based on angles formed by the face reconstruction where each face blends to an upwind scheme independently. Specifically we are concerned with two potentially problematic angles: 1) the angle between the face normal and the vector joining the cell center should be as close to zero as possible, and 2) the angle of a line that is tangent to a sphere centered on the skew-symmetric flux center (X_c) with a radius containing the face centroid, $r_f = |X_f - X_c|$ as illustrated in figure 5. The geometric upwinding is computed based on the maximum of these two angles, denoted as θ_{max} . Then the geometric based upwinding is computed using the relation

$$\alpha_{geom} = \chi^2 + \frac{\eta_2}{\eta_1} * (\chi^3 - 2\chi^2 + \chi); \chi = \min\{\eta_1(1 - \cos \theta_{max}), 1\}\tag{110}$$

where η_1 selects the angle at which full upwinding will occur, and η_2 controls how fast upwinding is added for small departures from ideal mesh quality. For the test cases studied it was found that accurate and robust results were obtained with, $\eta_1 = 2$, which provides full upwinding when $\theta_{max} = 60$ deg, while $\eta_2 = 1$ ensures stability on good quality meshes where small corrections are applied keeping $\alpha_{geom} \sim 1 - \cos \theta_{max}$ for small angles. In general, η_1 and η_2 are user adjustable parameters, but the settings described here are effective for a wide range of mesh configurations and appear to provide a reasonable compromise between robustness and low dissipation.

Since the central flux scheme cannot properly resolve shocks, a method must be used to identify regions of the flow that have shocks where it is necessary to switch to an upwind scheme. One of the more popular shock detectors used for this purpose is the method originally implemented by Ducros *et al.*[28] which detects shocks by comparing the magnitude of the divergence of velocity with the vorticity. Regions that are vorticity dominant (such as LES regions) switch to the low dissipation flux. However, this switch often becomes active in regions where an upwind scheme is not necessary. In addition, refining a region of smooth flow will not allow the code to switch to a lower dissipation flux. Ideally a switch should reduce the area where the upwind scheme is utilized with mesh refinement, and so a switch that considers local mesh resolution is desired. This switch should be based on the magnitude of divergence of the velocity field which gives a measure of the rate of compression in the flow, a mesh length scale, and some measure of the fluid response to compression, such as sound speed. Using these measures and dimensional




Figure 5: Geometry of flux stencil formation

analysis we arrive at the following

$$\alpha_{comp} = \left[\min \left(\eta_{comp} \frac{|\nabla \cdot \vec{u}| \delta x_{ref}}{a}, 1 \right) \right]^2 \quad (111)$$

where a is the local sound speed, and δx_{ref} is the cell reference length defined as the diameter of the smallest sphere that encloses all cell face centers. The variable η_{comp} is a user defined sensitivity parameter. The non-dimensional term is squared so that the transition to the central scheme will correspond with mesh refinement at a rate of $O(h^2)$ in smooth flow regions. For the studies that involve shocks presented here a value of $\eta_{comp} = 5$ is used. Note, with this method the value for α_{comp} will vanish with mesh refinement for any smooth solution, but will always be active at shocks. The final value for α_{diss} used by the flux is defined by the maximum of the left and right values for α_{comp} and the local face value for α_{geom} .

D.5 Time Integration

For the time integration scheme we utilize a Crank-Nicholson inspired scheme where the residual is evaluated on an interpolated half step. The interpolated half step, denoted by the primitive variable vector q^* , is computed according to the relation

$$q^* = [T^*, \vec{u}^*, p^*], \quad (112)$$

where

$$\begin{aligned} T^* &= \frac{2T^{n+1} T^n}{T^{n+1} + T^n} \\ \vec{u}^* &= \frac{\sqrt{\rho^{n+1}} \vec{u}^{n+1} + \sqrt{\rho^n} \vec{u}^n}{\sqrt{\rho^{n+1}} + \sqrt{\rho^n}} \\ p^* &= \frac{(1 + \epsilon)p^{n+1} + (1 - \epsilon)p^n}{2} \end{aligned} \quad (113)$$

This differs from the Subbareddy *et al.* [23] scheme in several important aspects. First the average is defined in terms of pressure and temperature to improve the performance of the algorithm when there is a more complex EoS as mentioned previously. Second, a harmonic averaging of temperature is utilized as this provides stability in high speed flow cases. The

energy consistent form of velocity averaging is retained. Note, the ϵ pressure biasing in the pressure averaging is used to provide the scheme with acoustic damping which is necessary when running at high Courant-Friedrich-Lewy (CFL) numbers. The effect of this term will be elucidated in later discussions. The time integration can then be expressed as the solution to the non-linear equation

$$\mathcal{L}(q^*) = \frac{\mathcal{V}}{\Delta t}(Q^{n+1} - Q^n) - R(q^*) = 0, \quad (114)$$

where $Q = [\rho, \rho\vec{u}, \rho e_T]$ is the conservative variable vector where the total energy is $e_T = e(p, T) + (\vec{u} \cdot \vec{u})/2$. The residual, R , results from the numerically integrated flux functions given by

$$R = \sum_{f=1}^{nf} \mathcal{A}_f (F_{v,f} - F_{i,f}) \quad (115)$$

where \mathcal{A}_f is the area of face f , and F_v and F_i are the viscous and inviscid fluxes respectively.

The solution is advanced in time by solving eq. (114) by a Newton iteration method which can be described as follows

$$\mathcal{L}'(q^{*,n+1,k})(q^{*,n+1,k+1} - q^{*,n+1,k}) = -\mathcal{L}(q^{*,n+1,k}), \quad (116)$$

for $k \geq 0$, where the Newton iteration is initialized using the previous time step value ($q^{*,n+1,k=0} = q^{*,n}$). To limit the size of the jacobian of operator \mathcal{L} , only the first order parts of the flux stencils are retained. For the inviscid flux jacobians we utilize approximate jacobians of the HLLC flux as described by Batten *et al.* [29] even when the skew-symmetric fluxes are utilized. These jacobian approximation errors have no effect on the solution once the Newton method is converged.

D.5.1 Local Time-Stepping Scheme

When running in steady timestepping mode the code utilizes a local timestep based on estimated changes in local pressure and temperature. This estimate is base on a first order explicit step whereby the estimated rate of change in pressure and temperature can be determined through solving the inverse of the jacobian of the conservative variables with respect to the primitive variables times the source terms formed from the inviscid and viscous fluxes. The timestep that will limit the change of these variables no more than the μ_{relax} factor is computed from these rates. If this timestep is smaller than the specified maximum timestep, then the smaller value will be used.

D.6 Scalar Transport Equations

The scalar transport equation in conservative form is written for a convected scalar, ϕ as

$$\frac{\partial \rho \phi}{\partial t} + \nabla \cdot (\rho u \phi) = \nabla \cdot (\lambda \nabla \phi) + S_\phi, \quad (117)$$

where ρ is the fluid density, u is the fluid velocity vector, λ is a diffusion coefficient, and S_ϕ is a prescribed source term. We note that by substituting $\phi = 1$ into the above equation one recovers the global continuity equation. We discretize the above equation in time and space to arrive at the expression

$$\begin{aligned} \frac{\mathcal{V}}{\Delta t} [(1 + \psi)(\phi^{n+1} \rho^{n+1} - \phi^n \rho^n) - \psi(\phi^n \rho^n - \phi^{n-1} \rho^{n-1})] = \\ - \sum \phi_u^{n+1} \dot{m}_f + \sum D_{\phi_f} + S_\phi, \end{aligned} \quad (118)$$

where ϕ_u is an upwinded extrapolation of ϕ based on the sign of the face mass flux \dot{m} and D_{ϕ_f} is the numerical diffusion flux for scalar ϕ . Note that the discrete continuity equation is given by

$$\frac{\mathcal{V}}{\Delta t} [(1 + \psi)(\rho^{n+1} - \rho^n) - \psi(\rho^n - \rho^{n-1})] = -\sum \dot{m}_f. \quad (119)$$

To improve the diagonal dominance of the scheme we employ a trick where we recognize that Eq. (119) is zero when the system of equation is solved, thus we multiply Eq. (119) by ϕ^{n+1} and subtract it from Eq. (118) to arrive at the following expression:

$$\begin{aligned} \frac{\mathcal{V}}{\Delta t} [(1 + \psi)(\phi^{n+1} - \phi^n)\rho^n - \psi(\phi^n\rho^n - \phi^{n-1}\rho^{n-1}) + \psi\phi^{n+1}(\rho^n - \rho^{n-1})] - \\ \left[\phi^{n+1} \sum \dot{m}_f - \sum \phi_u^{n+1} \dot{m}_f + \sum D_{\phi_f} + S_{\phi} \right] = L_{\phi}(\phi^{n+1}) = 0. \end{aligned} \quad (120)$$

The above operator, L_{ϕ} , is solved using a Newton method concurrently with the fluid equations.

References

- [1] E. Luke and T. George. Loci: A Rule-Based Framework for Parallel Multidisciplinary Simyulation Synthesis. *Journal of Functional Programming*, 14(03):477–502, 2005.
- [2] Y. Zhang and E. Luke. Concurrent Composition Using Loci. *Computing in Science and Engineering*, 11(3):27–35, May 2009.
- [3] Josh Barnes and Piet Hut. A hierarchical $O(N \log N)$ force-calculation algorithm. *Nature*, *Nature Publishing Group*, 324:446–449., December 1986.
- [4] L Greengard and V Rokhlin. A fast algorithm for particle simulations. *Journal of Computational Physics*, 73(2):325 – 348, 1987.
- [5] B.S. Baldwin and H. Lomax. Thin Layer Approximation and Algebraic Model for Separated Turbulent Flows. In *Proceedings of the AIAA 16th Aerospace Sciences Meeting*. AIAA, January 1978. AIAA-78-257.
- [6] P. R. Spalart and S. R. Allmaras. A One-equation Turbulence Model for Aerodynamic Flows. *AIAA-92-0439*, 1992.
- [7] F. R. Menter. Influence of Freestream Values on $k - \omega$ Turbulence Model Predictions. *AIAA Journal*, 30(6):1657–1659, June 1992.
- [8] D. C. Wilcox. Reassessment of the Scale-Determining Equation for Advanced Turbulence Models. *AIAA Journal*, 26(11):1299–1310, 1988.
- [9] D. C. Wilcox. A Half Century Review of the $k - \omega$ Model. Technical report, AIAA, 1991. AIAA 91-1784.
- [10] F. R. Menter. Two-Equation Eddy-Viscosity Turbulence Models for Engineering Applications. *AIAA Journal*, 32(8):1598–1605, August 1994.
- [11] D. C. Wilcox. *Turbulence Modeling for CFD*. DCW Industries, 1998.
- [12] D. A. Johnson and L. S. King. Mathematically Simple Turbulence Closure Model for Attached and Separated Turbulent Boundary Layers. *AIAA Journal*, 23(11):1684–1692, November 1985.

- [13] D. C. Wilcox. Formulation of the $k-\omega$ Turbulence Model Revisited. *AIAA J.*, 46(11):2823–2838, 2008.
- [14] R.H. Nichols and C.C. Nelson. Application of Hybrid RANS/LES Turbulence Models. Technical report, AIAA, 2003. AIAA 2003-0083.
- [15] M. Samimy and G. S. Elliott. Effects of Compressibility on the Charactersitics of Free Shear Layers. *AIAA Journal*, 28(3):439–445, March 1990.
- [16] S. Sarkar and B. Lakshmanan. Application of a Reynolds Stress Turbulence Model to the Compressible Shear Layer. *AIAA Journal*, 29(5):743–749, May 1991.
- [17] S. Bhushan and D. K. Walters. A dynamic hybrid Reynolds-averaged Navier Stokes–Large eddy simulation modeling framework. *Physics of Fluids*, 24(1):015103, 2012.
- [18] D. K. Walters, S. Bhushan, M. F. Alam, and D. S. Thompson. Investigation of a Dynamic Hybrid RANS/LES Modelling Methodology for Finite-Volume CFD Simulations. *Flow, Turbulence and Combustion*, 91(3):643–667, 2013.
- [19] R. Koomullil and B. Soni. Flow Simulation Using Generalized Static and Dynamic Grids. *AIAA Journal*, 37(12):1551–1557, December 1999.
- [20] A. G. Hansen. Generalized Control Volume Analyses with Application to the Basic Laws of Mechanics and Thermodynamics. *Bulletin of Mechanical Engineering Education*, 4:161–168, 1965.
- [21] P. D. Thomas and C. K. Lombard. Geometric Conservation Law and Its Application to Flow Computations on Moving Grids. *AIAA Journal*, 17(10):1030–1037, 1978.
- [22] E.F. Toro, M. Spruce, and M. Speares. Restoration of the contact surface in the HLL-Riemann solver. *Shock Waves*, 4:25–34, 1994.
- [23] P.K. Subbareddy and G.V. Candler. A fully discrete, kinetic energy consistent finite-volume scheme for compressible flows. *Journal of Computational Physics*, 228(5):1347–1364, 2009.
- [24] S. Pirozzoli. Generalized conservative approximations of split convective derivative operators. *Journal of Computational Physics*, 229(19):7180–7190, 2010.
- [25] S. Pirozzoli. Stabilized non-dissipative approximations of Euler equations in generalized curvilinear coordinates. *Journal of Computational Physics*, 230(8):2997–3014, 2011.
- [26] M. S. Darwish and F. Moukalled. TVD schemes for unstructured grids. *International Journal of Heat and Mass Transfer*, 46:599–611, 2003.
- [27] B. Thornber, D. Drikakis, R.J.R. Williams, and D. Youngs. An improved reconstruction method for compressible flows wiht low Mach number features. *Journal of Computational Physics*, 227:4853–4872, 2008.
- [28] F. Ducros, F. Laporte, T. Souleres, V. Guinot, P. Moinat, and B. Caruelle. High-order fluxes for conservative skew-symmetric like schemes in structured meshes: Application to compressible flows. *Journal of Computational Physics*, 161:114–139, 2000.
- [29] P. Batten, N. Clarke, C. Lambert, and D.M. Causon. On the Choice of Wavespeeds for the HLLC Riemann Solver. *SIAM Journal on Scientific Computing*, 18(6):1553–1570, 2006.